

NGINX JS DOM — Developer Manual

Contents

1. Introduction	4
1.1. What NGINX JS DOM Is	4
1.2. Intended Audience	5
1.3. A Map of the Moving Parts	5
2. Architecture	6
2.1. Lifecycle Overview	6
2.2. Process and Thread Topology	7
2.3. COM Ownership and Mutation Scope	8
3. Quick Start	8
3.1. Hello World — Minimal (No http{} block)	9
3.2. Hello World — Classic (with http{} block)	9
3.3. Reading Request Data	10
3.4. Async Fetch and Transform	11
3.5. Live Upstream Rebalancing	12
4. Request Handlers	12
4.1. Installing a Handler	12
4.2. NginxRequest Reference	13
4.2.1. Identity and Network	13
4.2.2. Headers	13
4.2.3. Body	13
4.2.4. Cookies and Variables	14
4.2.5. Location and Server Context	14
4.2.6. Response	14
4.2.7. Async Utilities	14
4.2.8. Logging	15
4.3. Async Handler Lifecycle	15
4.4. Streaming Responses	16
4.5. Subrequest Fan-Out	16
4.6. Header and Body Filters	16
5. Configuration Object Model (COM)	17
5.1. Overview	17
5.2. NginxServer	18
5.3. NginxLocation	18
5.4. NginxUpstream and Peers	19
5.5. Iterating the Full Config Tree	20
6. Dynamic Reconfiguration	20

6.1. Scope 1 — This Worker Only	20
6.2. Scope 2 — All Workers Atomically	20
6.3. Scope 3 — Add / Remove Servers and Locations	21
6.4. Scope 4 — Per-Request Config Snapshot	22
6.5. Hot SSL Certificate Swap	23
7. Pre-Fork Initialization	23
7.1. <code>js_preprocess</code> — Config Injection During Parse	23
7.1.1. Example: Server Blocks from a JSON File	24
7.1.2. Example: Conditional SSL Config	24
7.1.3. Example: Fetching Config from a Network Service	25
7.1.4. Multiple <code>js_preprocess</code> Directives	26
7.2. <code>js_init_http</code> — Post-Parse Server Injection	26
8. Workers (JS Web Workers)	27
8.1. Basic Worker Pattern	27
8.2. Transferring ArrayBuffers (Zero Copy)	28
8.3. Worker <code>std</code> and <code>os</code> Modules	29
9. SharedWorkers	29
9.1. Concept and Use Cases	29
9.2. Port-Based Pattern (Explicit <code>onconnect</code>)	29
9.3. Broadcast Pattern (Global <code>onmessage</code> + <code>clients</code>)	31
9.4. Dynamic SharedWorker Creation	32
9.5. SharedWorker Flow Chart	32
10. SharedArrayBuffers and Atomics	33
10.1. What They Are	33
10.2. Pre-Fork SABs (Visible to All Workers)	33
10.3. Post-Fork SABs (Worker-Local, Passed via Channels)	33
10.4. Atomics.wait / Atomics.notify	34
11. Accept Control and Cross-Worker Coordination	35
11.1. The Problem	35
11.2. Per-Worker: <code>suspendAcceptance</code> / <code>resumeAcceptance</code>	35
11.3. Cross-Worker: <code>suspendAllWorkers</code> / <code>resumeAllWorkers</code>	36
11.4. SharedWorker Broadcast + Accept Control Pattern	37
11.5. Backlog Behaviour	38
12. Snapshots	38
12.1. Basic Snapshot and Rollback	38
12.2. Sub-Object Snapshots	39
12.3. Read / Write Mode Interaction with Snapshots	39
13. Stream (TCP/UDP) Support	40
14. Bootstrap Functions (<code>nginx.createHttp</code> / <code>nginx.createStream</code>)	41
14.1. <code>nginx.createHttp([directives])</code>	41
14.2. <code>nginx.createStream([directives])</code>	42

14.3. When to Use Bootstrap Functions	42
15. Master/Worker Scripting	42
15.1. Master Lifecycle Events (<code>nginx.on</code>)	43
15.2. Master → Worker Messages	44
15.3. Worker → Master Messages	44
15.4. <code>nginx.workerIdx</code>	45
16. <code>nginx.repl</code> — Interactive JavaScript REPL	45
16.1. How It Works	45
16.2. REPL API	45
16.3. Wire Protocol	46
16.4. Embedding the REPL	46
16.5. The <code>nginx-repl</code> Tool	47
17. Applications	48
17.1. <code>admin_snapshot_rollback</code>	48
17.1.1. Concept	48
17.1.2. Admin HTTP API	48
17.1.3. JavaScript API (<code>nginx.admin</code>)	48
17.1.4. <code>location.clearHandler()</code>	49
17.1.5. Running the Demo	49
18. External Config Sources and Full Reload	49
18.1. Pattern A — File at Startup (<code>js_preprocess</code>)	50
18.2. Pattern B — Network at Startup (<code>js_preprocess</code> + <code>std.urlGet</code>)	50
18.3. Pattern C — Live Reload via Handler (Zero Downtime)	50
18.4. Decision Guide	52
19. Third-Party Module Integration	52
19.1. Overview	52
19.2. Step 1 — Identify the Config Structs	52
19.3. Step 2 — Write the COM Bridge File	53
19.4. Step 3 — Register the Bridge in the COM Initialiser	55
19.5. Step 4 — Use from JavaScript	55
19.6. Location-Level Config	56
20. Resource Protection	56
20.1. Worker Memory Limit	56
20.2. Worker Request Timeout	57
20.3. Error Isolation	57
20.4. What JS Cannot Do	57
21. Reference	58
21.1. Directives	58
21.2. <code>nginx</code> Global Object	58
21.3. <code>nginx.http</code> Object	60
21.4. <code>NginxRequest</code> Methods and Properties	60

21.5. Worker / SharedWorker API	60
21.6. SharedArrayBuffer and Atomics Support Matrix	61
21.7. Log Level Constants	61
22. Appendix — Flow Charts	61
22.1. A: Simple Synchronous Handler	62
22.2. B: Async Handler with Worker Offload	62
22.3. C: Cross-Worker COM Mutation (suspendAllWorkers)	62
22.4. D: SharedWorker State Coordination	62
22.5. E: SAB Lock-Free Counter (Pre-Fork)	62
22.6. F: Pre-Fork Config Generation	63
22.7. G: Live Reload from Config Service	63
22.8. H: Complete nginx.* Object Tree	63
22.8.1. NginxServer properties	65
22.8.2. NginxLocation properties	65
22.8.3. NginxUpstream, NginxPeer, NginxRRPeer	68
22.8.4. NginxSocket and NginxHttpListener / NginxStreamListener	68
22.8.5. NginxStreamServer and NginxStreamProxy	69
22.8.6. NginxRequest and NginxStreamSession	69
22.8.7. Worker and SharedWorker	71
22.8.8. NginxSSL and NginxStreamSSL	71
22.9. I: Full Nginx JS Class Reference	71
22.9.1. Class Index	71
22.9.2. Module COM Class Properties	74

1. Introduction

1.1. What NGINX JS DOM Is

NGINX is deliberately static: its configuration is read once at startup (or reload) and compiled into kernel-level data structures. This rigidity is a feature — it makes nginx predictable, fast, and easy to reason about — but it creates friction whenever runtime adaptability is required: upstream weights that need to track backend health scores, virtual hosts generated from a tenant database, SSL certificates that rotate on a schedule, or request routing rules that depend on live feature-flag state.

NGINX JS DOM embeds [QuickJS](#) (a lightweight [ES2020+](#) JavaScript engine) directly into nginx and exposes the entire nginx configuration as a structured JavaScript object tree called the **Configuration Object Model (COM)**. Through the COM, JavaScript can:

- **Read** the full live configuration of every module in every server and location.
- **Write** configuration properties — upstream peer weights, SSL certificates, location directives, gzip settings, and much more — with immediate effect, either for the current worker or atomically for all workers at once.

- **Generate** server and location blocks at parse time, before nginx has even finished reading its configuration file, and at run-time, when traffic events are processed by nginx workers in their nginx event loops.
- **Handle requests** in JavaScript, using the same event-driven async model as nginx itself.
- **Offload** CPU-intensive work to background threads (Workers) without blocking the nginx event loop.
- **Maintain shared mutable state** across all nginx worker processes using SharedWorkers and SharedArrayBuffers.

The result is an nginx that can be programmed rather than merely configured — without restarting, without patching C code, and without sacrificing the performance or stability of the core server.

1.2. Intended Audience

This manual assumes:

- Familiarity with nginx configuration syntax and core concepts (workers, upstreams, locations, modules). See the [NGINX Beginner's Guide](#) and the [full module reference](#) if you need a refresher.
- Working knowledge of modern JavaScript (ES2020): `async/await`, Promises, `import/export`, `SharedArrayBuffer`, `Atomics`. See the [ECMAScript 2020 specification](#) or any ES2020 reference.

You do **not** need to know C or QuickJS internals to use the JS DOM API.

1.3. A Map of the Moving Parts

The sections that follow cover these components:

Component	What it gives you
<code>js_preprocess</code>	Runs JavaScript during config parsing; the script writes config text back into the parser. Use it to generate server/location blocks from files, environment variables, or external sources.
<code>js_init_http</code>	Runs JavaScript after the <code>http {}</code> block is parsed but before directives are merged. Can read already-parsed servers and inject new ones.
<code>js_source</code>	Evaluates JavaScript in the master process after all config has been parsed. This is where you install request handlers, set up Workers and SharedWorkers, and perform any pre-fork initialisation.
Request handlers	JavaScript <code>async</code> functions assigned to <code>location.handler</code> . Each request runs through the nginx event loop; <code>await</code> suspends the handler without blocking any other request.
COM	The <code>nginx</code> global object tree. Read live config values; write mutations that take effect immediately.

Component	What it gives you
Workers	Standard Web Worker API. A per-nginx-worker background thread for CPU-bound work. Does not block the event loop.
SharedWorkers	A single long-lived thread in the master process, connected to every nginx worker via UNIX sockets. The canonical place for mutable state that must be consistent across all workers.
SharedArrayBuffers + Atomics	Binary shared-memory regions (MDN reference) accessible from any thread with zero copy. Combined with Atomics , they enable lock-free counters, mutexes, and condition variables.
Bootstrap functions	<code>nginx.createHttp()</code> / <code>nginx.createStream()</code> — instantiate the HTTP or stream module entirely from JavaScript, without an <code>http {}</code> or <code>stream {}</code> block in <code>nginx.conf</code> . Combine with <code>nginx.http.addServer()</code> and <code>nginx.createSocket()</code> for a fully JS-defined server topology.
Master/worker scripting	<code>nginx.on(event, fn)</code> hooks into master lifecycle events (worker spawn, reload, SIGUSR1, ...). <code>nginx.sendToWorker()</code> / <code>nginx.broadcastToWorkers()</code> push messages from master to workers; <code>nginx.sendToMaster()</code> sends from a worker back to the master.
<code>nginx.repl</code>	An interactive JavaScript REPL that evaluates expressions inside a running nginx worker over a plain HTTP upgrade connection. Useful for live introspection, ad-hoc mutations, and debugging.
<code>location.clearHandler()</code>	Removes a previously installed JS content handler and restores the original nginx C handler for that location. Used by snapshot/rollback to undo handler changes atomically.
Accept control	<code>nginx.suspendAcceptance()</code> / <code>nginx.suspendAllWorkers()</code> — pause connection acceptance during a mutation window, then resume.
Resource limits	<code>nginx.workerMemoryLimit</code> and <code>nginx.workerRequestTimeout</code> cap heap usage and CPU time per request; JS exceptions are caught; the worker survives.


2. Architecture

2.1. Lifecycle Overview


```


nginx start
|
|— Config file parsing begins (ngx_conf_parse)
|   |
|   |— js_preprocess /gen.js — STAGE 1: runs immediately, mid-parse
|   |   □ Ephemeral JS runtime (freed after script returns)
|   |   □ config.write(text) injects text back into the parser
|   |   □ nginx.http, nginx.stream, etc. do NOT exist yet
|   |
|   |— ... rest of nginx.conf parsed (servers, locations, upstreams) ...

```



2.2. Process and Thread Topology





Key properties of this topology:

- Each nginx **worker process** has its own independent JS runtime. QuickJS is not thread-safe; no runtime is ever shared between processes.
- Workers inherit the master's runtime via copy-on-write after `fork()`. Scripts do not need to be re-evaluated; all handler functions, SharedWorker references, and global variables are ready immediately.
- **SharedWorkers** run in the master process as pthreads, outside any worker. Each worker communicates with a SharedWorker over a dedicated UNIX socket channel. This is the only channel through which workers can share mutable JS state.
- **JS Workers** (lowercase) are pthreads spawned **inside** a worker process. They exist to offload work from the event loop and are private to their nginx worker.

2.3. COM Ownership and Mutation Scope

All COM mutations go through one of three scopes:

Scope	Mechanism	Visible to
Per-request snapshot	<code>r.location.setWriteMode("local")</code> — writes a per-request copy of the location config	This request only; reversed when the request pool is freed
This worker only	Direct COM write (default)	All subsequent requests on this worker process
All workers atomically	<code>await nginx.suspendAllWorkers()</code> → mutate → <code>await nginx.resumeAllWorkers()</code>	All workers, once they resume

3. Quick Start

This section goes from "hello world" to "live upstream rebalancing" in four steps, each building on the previous one.

3.1. Hello World — Minimal (No http{} block)

The most compact nginx JS DOM server needs only a `js_source` directive and an `events {}` block. Everything else — the HTTP module, socket, virtual server, and location handler — is created from JavaScript at init time:

```
# nginx.conf
js_source /etc/nginx/js/hello_world.js;

events { }
```

```
// hello_world.js – evaluated once in the master process after config parsing.
// Creates the entire server topology from JS; no http{} / server{} / location{}
// blocks are needed in nginx.conf.
```

```
(function () {

    // 1. Bootstrap the HTTP module (equivalent to an empty http{} block).
    //    Pass any http-level directives as a string argument.
    nginx.createHttp('access_log /var/log/nginx/access.log;');

    // 2. Create a new virtual server.
    var srv = nginx.http.addServer('hello-world.local');

    // 3. Add a location and install the JS content handler inline.
    srv.addLocation('/hello').handler = function (req) {
        req.respond(200, { 'content-type': 'text/plain' },
            'JS COM: Hello World!\n');
    };

    // 4. Bind a TCP socket, attach it to the HTTP module,
    //    and register the server on that listener.
    var sock = nginx.createSocket('0.0.0.0:8080');
    nginx.http.attach(sock).addServer(srv);

})();
```

3.2. Hello World — Classic (with http{} block)

When `http {}` is already present in `nginx.conf`, just find the location placeholder and assign a handler function:

```
# nginx.conf
js_source /etc/nginx/js/init.js;

events { }
```

```
http {
    server {
        listen 80;
        location /hello { }    # placeholder – handler installed from JS
    }
}
```

```
// /etc/nginx/js/init.js
// This file is evaluated once in the master process after config parsing.
// All location.handler assignments made here are inherited by every worker.

const loc = nginx.http.servers[0].locations.find(l => l.path === '/hello');

loc.handler = (req) => {
    // req is a NginxRequest object. text() is a convenience shorthand
    // for respond(200, {'Content-Type': 'text/plain'}, body).
    req.text('Hello from JS DOM!\n');
};
```

3.3. Reading Request Data

```
// Read everything you might want from an incoming request.

loc.handler = async (req) => {
    // Basic identity
    const method = req.method;           // "GET", "POST", ...
    const uri    = req.uri;               // "/hello"
    const args   = req.args;              // raw query string: "a=1&b=2"
    const params = req.queryParams;       // {a: "1", b: "2"}
    const scheme = req.scheme;            // "http" or "https"
    const host   = req.host;              // "example.com"

    // Network info
    const client = req.remoteAddr;        // "203.0.113.42"
    const port   = req.remotePort;        // 54321
    const server = req.serverAddr;        // "192.0.2.1"

    // Headers (read-only object; keys are lowercased)
    const ua     = req.headers['user-agent'];
    const ct     = req.contentType;       // shorthand for content-type header

    // Cookies: {name: "value", ...}
    const session = req.cookies['session_id'];

    // nginx variables – readable and writable via req.variables
    const xff     = req.variables['http_x_forwarded_for'];
    req.variables['my_custom_var'] = 'set-from-js';
}
```

```
// Request body (requires await – nginx reads it asynchronously)
const body    = await req.readBody(); // returns string

req.json({
  method, uri, client, ua, session,
  bodyLength: body.length
});
};
```

3.4. Async Fetch and Transform

```
// Non-blocking upstream call. While this handler awaits, the nginx event
// loop freely serves other requests on this worker.

loc.handler = async (req) => {
  // req.fetch() performs an HTTP subrequest through nginx's upstream
  // machinery. It returns a Promise that resolves with a Response object.
  const res = await req.fetch('http://api.internal/v1/prices');
  const data = await res.json();

  // Apply transformation
  const adjusted = data.items.map(item => ({
    ...item,
    price: (item.price * 1.10).toFixed(2) // add 10% margin
  }));

  req.json({ items: adjusted });
};
```



3.5. Live Upstream Rebalancing

```
// An admin endpoint that atomically updates upstream peer weights across
// all workers – no nginx restart, no dropped connections.

const adminLoc = nginx.http.servers
  .find(s => s.names.includes('admin.internal'))
  .locations.find(l => l.path === '/upstream/rebalance');

adminLoc.handler = async (req) => {
  const body    = await req.readBody();
  const weights = JSON.parse(body);    // { "192.0.2.10:8080": 5, "192.0.2.11:8080":
10 }

  // Step 1: pause all workers so the change is seen atomically.
  await nginx.suspendAllWorkers();

  // Step 2: apply weights.
  const peers = nginx.http.upstreams
    .find(u => u.name === 'backend').peers;

  for (const peer of peers) {
    if (weights[peer.addr] !== undefined) {
      peer.weight = weights[peer.addr];
    }
  }

  // Step 3: resume – all workers pick up the new weights immediately.
  await nginx.resumeAllWorkers();

  req.json({ applied: true, peers: peers.map(p => ({addr: p.addr, weight:
p.weight})) });
};
```

4. Request Handlers

4.1. Installing a Handler

A location handler is a JavaScript function assigned to `location.handler` during `js_source` evaluation (`init_conf` time):

```
nginx.http.servers[0].locations
  .find(l => l.path === '/api')
  .handler = async (req) => { ... };
```

The function receives one argument — a `NginxRequest` object — and may be synchronous or `async`. A

location that has no JS handler falls through to nginx's normal content processing (proxy, static files, etc.).

4.2. NginxRequest Reference

4.2.1. Identity and Network

Property	Type	Description
<code>method</code>	string	HTTP method: "GET", "POST", "DELETE", ...
<code>uri</code>	string	Request URI without query string
<code>args</code>	string	Raw query string ("a=1&b=foo")
<code>queryParams</code>	object	Query string parsed to key → value pairs
<code>scheme</code>	string	"http" or "https"
<code>host</code>	string	Value of the <code>Host</code> header
<code>remoteAddr</code>	string	Client IP address
<code>remotePort</code>	number	Client TCP port
<code>serverAddr</code>	string	Server-side IP the connection arrived on
<code>serverPort</code>	number	Server-side port
<code>httpVersion</code>	string	"1.0", "1.1", or "2.0"
<code>keepalive</code>	boolean	Whether the connection is keep-alive
<code>isInternal</code>	boolean	<code>true</code> for subrequests
<code>startTime</code>	number	Request start time (milliseconds, Unix epoch)
<code>requestLength</code>	number	Total bytes received (headers + body)
<code>connection</code>	number	nginx connection serial number

4.2.2. Headers

Property / Method	Type	Description
<code>headers</code>	object	Request headers; keys are lowercased
<code>contentType</code>	string	Shorthand for <code>headers['content-type']</code>
<code>contentLength</code>	number	Parsed <code>Content-Length</code> value
<code>getHeader(name)</code>	string	Get a response header already set
<code>setHeader(name, v)</code>	—	Set a response header
<code>removeHeader(name)</code>	—	Remove a response header

4.2.3. Body

Property / Method	Type	Description
<code>body</code>	string null	Request body if already buffered; <code>null</code> otherwise
<code>readBody()</code>	Promise<string>	Read the full request body (async)

4.2.4. Cookies and Variables

Property / Method	Type	Description
<code>cookies</code>	object	Cookie jar: <code>{name: "value", ...}</code>
<code>variables</code>	object	Exotic r/w proxy over nginx variables
<code>getVar(name)</code>	string	Read one nginx variable
<code>setVar(name, val)</code>	—	Set one nginx variable

4.2.5. Location and Server Context

Property	Type	Description
<code>location</code>	NginxLocation	The matched location's COM object (mutable within this request)
<code>upstream</code>	object null	Upstream peer info if request was proxied

4.2.6. Response

Method	Description
<code>respond(status, headers, body)</code>	Low-level response; sends status + headers + body string in one call
<code>json(data)</code>	Send JSON: status 200, <code>Content-Type: application/json</code>
<code>text(body)</code>	Send plain text: status 200, <code>Content-Type: text/plain</code>
<code>html(body)</code>	Send HTML: status 200, <code>Content-Type: text/html</code>
<code>redirect(url)</code>	Send HTTP 302 redirect
<code>sendfile(path)</code>	Serve a file from the filesystem using <code>sendfile(2)</code>
<code>writeHead(status, headers)</code>	Start a streaming response (headers only)
<code>write(chunk)</code>	Append a chunk to a streaming response
<code>finish()</code>	End a streaming response
<code>statusCode</code> (r/w)	Get or set the response status code

4.2.7. Async Utilities

Method	Description
<code>sleep(ms)</code>	Pause the handler for <code>ms</code> milliseconds (non-blocking)

Method	Description
<code>fetch(url [, opts])</code>	Make a non-blocking upstream HTTP request; returns <code>Promise<Response></code>
<code>subrequest(uri)</code>	Issue an nginx internal subrequest; returns <code>Promise<Response></code>

4.2.8. Logging

```
req.log(nginx.INFO, 'Processing request from ' + req.remoteAddr);
// Log levels: nginx.DEBUG, nginx.INFO, nginx.NOTICE, nginx.WARN,
//             nginx.ERR, nginx.CRIT, nginx.ALERT, nginx.EMERG
```

4.3. Async Handler Lifecycle

```
nginx worker event loop
├── accept() → new connection
├── HTTP request parsed
├── location matched → location.handler(req) called
│   ├── [sync path] req.json(data)
│   │   └── response buffered → written to socket → done
│   └── [async path] await something
│       ├── Handler suspended. Worker event loop continues:
│       │   ├── accepts new connections
│       │   ├── runs other handlers
│       │   └── processes timers and upstream responses
│       └── event fires (timer / upstream / Worker message)
│           ├── Promise resolves; handler resumes from the await point
│           └── req.json(data) → done
└── (next event)
```



Multiple async handlers can be suspended concurrently on the same worker. Each nginx worker runs a single-threaded event loop; async handlers cooperate by yielding at every `await` point, allowing other requests to be accepted and served in between. Use Workers or SharedWorkers when you need true CPU parallelism or need to share mutable state across workers.

4.4. Streaming Responses

```
loc.handler = async (req) => {
  // Send status + headers, then stream the body in chunks.
  req.writeHead(200, {'Content-Type': 'text/event-stream',
                     'Cache-Control': 'no-cache'});

  for (let i = 0; i < 5; i++) {
    req.write(`data: event ${i}\n\n`);
    await req.sleep(200);      // 200 ms between events, non-blocking
  }

  req.finish();
};
```

4.5. Subrequest Fan-Out

```
// Fire three upstreams in parallel and merge results.
loc.handler = async (req) => {
  const [a, b, c] = await Promise.all([
    req.subrequest('/internal/prices'),
    req.subrequest('/internal/stock'),
    req.subrequest('/internal/ratings'),
  ]);

  req.json({
    prices: await a.json(),
    stock:  await b.json(),
    ratings: await c.json(),
  });
};
```

4.6. Header and Body Filters

Filters are JavaScript functions that run in the nginx filter chain. They are installed on a location object:

```
// Header filter – modify response headers before they leave nginx
loc.headerFilter = (req, headers) => {
  headers['X-Powered-By'] = 'NGINX JS DOM';
  delete headers['Server'];
  return headers;
};

// Body filter – transform response body chunks
loc.bodyFilter = (req, chunk, isLast) => {
```

```
// Replace "foo" with "bar" in every body chunk
return chunk.replace(/foo/g, 'bar');
};
```


Filters run in priority order (lower number = earlier in chain, default priority 50).

5. Configuration Object Model (COM)

5.1. Overview

The COM is a live JavaScript object tree rooted at the global `nginx` object. Properties read from the tree return current nginx struct values; writes (on mutable properties) take immediate effect.

```
nginx
├── version          - string  "1.27.x"
├── cpuCount         - number  (number of CPU cores)
├── workerMemoryLimit - number  (bytes; 0 = unlimited)  ← r/w
├── workerRequestTimeout - number (ms; 0 = unlimited)  ← r/w
├── log(level, msg)   - write to nginx error log
├── broadcast(msg)    - send msg to all workers' nginx.onbroadcast handler
├── suspendAcceptance() - per-worker accept pause (sync)
├── resumeAcceptance() - per-worker accept resume (sync)
├── suspendAllWorkers() - cross-worker quiesce (async, returns Promise)
├── resumeAllWorkers() - cross-worker resume (async, returns Promise)
├── cycle
│   ├── connections    - number (active connections)
│   ├── freeConnections - number
│   ├── openFilesN     - number
│   └── sockets[]      - listener socket objects
├── events
│   ├── workerConnections - number ← r/w
│   ├── multiAccept       - boolean ← r/w
│   └── acceptMutex       - boolean ← r/w
├── http
│   ├── servers[]       - array of NginxServer (see below)
│   ├── upstreams[]     - array of NginxUpstream
│   ├── variables{}     - nginx HTTP variables
│   ├── serverNamesHashMaxSize ← r/w
│   ├── serverNamesHashBucketSize ← r/w
│   ├── variablesHashMaxSize ← r/w
│   ├── variablesHashBucketSize ← r/w
│   └── addServer(name)  - create a new virtual server at runtime
```



5.2. NginxServer

Each element of `nginx.http.servers[]` represents a `server {}` block.

```

const srv = nginx.http.servers[0];

// Identification
srv.name          // primary server_name
srv.names         // array of all server_names
srv.listen        // array of listen address strings

// Child collections
srv.locations     // array of NginxLocation
srv.upstreams     // upstreams referenced from this server (if any)

// SSL (if ssl module is compiled in)
srv.ssl.cert      // path to certificate ← r/w
srv.ssl.key       // path to private key ← r/w
srv.ssl.protocols // "TLSv1.2 TLSv1.3" ← r/w
srv.ssl.ciphers   // cipher list ← r/w
srv.ssl.clientCert // path to CA cert for client auth ← r/w

// Dynamic location management
const loc = srv.addLocation('/new-path'); // returns NginxLocation
const loc = srv.addLocation('/api', { template: '/existing-path' }); // copy settings
                                                                    // from existing
loc

// Stream handler (stream {} servers only)
srv.handler = async (session) => { ... };

```

5.3. NginxLocation

```

const loc = nginx.http.servers[0].locations.find(l => l.path === '/api');

// Identification
loc.path      // "/api"
loc.exact     // true for "= /api" locations

```

```

loc.regex          // true for "~ /pattern" locations

// Handler
loc.handler = async (req) => { ... };

// Common directives (r/w)
loc.root           // document root path
loc.alias           // alias path
loc.index           // index file list
loc.tryFiles        // try_files argument list
loc.sendfile        // boolean
loc.zip             // boolean
loc.zipMinLength    // number

// Proxy
loc.proxy.pass       // upstream URL ← r/w
loc.proxy.setHeader  // {name: value} pairs ← r/w
loc.proxy.readTimeout // ms ← r/w

// Rate limiting, access, auth, headers – all module configs accessible
// (see §14 Reference for full list)

// Per-request write scoping
loc.setWriteMode("global") // write to global config (all future requests)
loc.setWriteMode("local")  // write only to a per-request snapshot
loc.setWriteMode("both")   // write to both (default)
loc.setReadMode("global")  // reads come from global config
loc.setReadMode("local")   // reads come from the snapshot (default)

```

5.4. NginxUpstream and Peers

```

const up = nginx.http.upstreams.find(u => u.name === 'backend');

up.name          // "backend"
up.peers          // array of NginxRRPeer

const p = up.peers[0];
p.addr           // "192.0.2.10:8080"
p.weight          // number ← r/w (takes a write lock internally)
p.maxFails        // number ← r/w
p.failTimeout     // ms ← r/w
p.down            // boolean ← r/w (marks peer as permanently down)
p.backup          // boolean (read-only)
p.conns           // number (active connections, live stat)
p.fails           // number (fail counter, live stat)

```

5.5. Iterating the Full Config Tree

```
// Dump the entire server/location tree to the error log.
for (const srv of nginx.http.servers) {
  nginx.log(nginx.INFO, `Server: ${srv.names.join(', ')}');
  for (const loc of srv.locations) {
    nginx.log(nginx.INFO, ` ${loc.path} handler=${loc.handler ? 'JS' :
'nginx'}`);
  }
}

// Find all locations that use a specific upstream
const byUpstream = nginx.http.servers
  .flatMap(s => s.locations)
  .filter(l => l.proxy?.pass?.includes('backend'));
```

6. Dynamic Reconfiguration

6.1. Scope 1 — This Worker Only

The simplest mutations happen directly in a handler with no coordination. The change takes effect immediately for all subsequent requests on this worker. Other workers retain their own CoW copy unchanged.

```
// Use-case: A/B flag toggle via admin endpoint, applied to this worker.
// Fine for experiments where eventual consistency across workers is acceptable.

nginx.http.servers[0].locations
  .find(l => l.path === '/admin/toggle-gzip')
  .handler = (req) => {
    // Toggle gzip for the /api location on THIS worker only
    const apiLoc = nginx.http.servers[0].locations
      .find(l => l.path === '/api');
    apiLoc.gzip = !apiLoc.gzip;
    req.text(`gzip is now ${apiLoc.gzip}`);
  };
```

6.2. Scope 2 — All Workers Atomically

For changes that must be consistent across every worker simultaneously, use `suspendAllWorkers` / `resumeAllWorkers`.

```
Worker A (handler)
  |
  └─ await suspendAllWorkers()
```

```

    └─ □ sends coordination signal to Workers B, C, ...
        Each worker finishes its current request, then:
        □ stops calling accept()
        □ sends ACK
    └─ resolves when ALL workers have ACK'd

└─ [mutation window]
    □ Direct COM writes apply to Worker A's copy
    □ SharedWorker broadcast distributes the same data to B, C, ...
      (each worker applies it from its own onmessage handler)

└─ await resumeAllWorkers()
    └─ □ all workers call ngx_enable_accept_events() and resume

└─ respond

```

```

adminLoc.handler = async (req) => {
  const newPeers = JSON.parse(await req.readBody());

  await ngx.suspendAllWorkers();

  // Apply to this worker
  const peers = ngx.http.upstreams[0].peers;
  newPeers.forEach((p, i) => {
    if (peers[i]) peers[i].weight = p.weight;
  });

  await ngx.resumeAllWorkers();

  req.json({ ok: true });
};

```

6.3. Scope 3 — Add / Remove Servers and Locations

`addServer`, `removeServer`, `addLocation` enable structural mutations — adding or removing routing entries at runtime.

```

// Tenant onboarding: create a new vhost live, no restart.

ngx.http.servers
  .find(s => s.names.includes('admin.internal'))
  .locations.find(l => l.path === '/tenant/create')
  .handler = async (req) => {
    const body = JSON.parse(await req.readBody());
    const domain = body.domain;

```

```

    await nginx.suspendAllWorkers();

    // Create the new server
    const srv = nginx.http.addServer(domain);
    srv.listen.push('80');

    // Add a root location with a JS handler
    const loc = srv.addLocation('/');
    loc.handler = (r) => {
        r.text(`Welcome to ${domain}!\n`);
    };

    // Rebuild the server_name dispatch hash so the new Host: domain routes
    // correctly to our new server.
    nginx.http.rebuildVhostDispatch();

    await nginx.resumeAllWorkers();

    req.json({ created: domain });
};

```

```

// Remove a tenant server (e.g. account deleted).
nginx.http.removeServer('ex-tenant.example.com');
nginx.http.rebuildVhostDispatch();

```



`rebuildVhostDispatch()` must be called after any change to server names. Without it the Host-based routing hash is stale and requests will be routed to the wrong server.

6.4. Scope 4 — Per-Request Config Snapshot

The `location.setWriteMode("local")` API creates a copy of the location's config struct scoped to the current request. The global struct is untouched; other concurrent requests see the original values.

```

loc.handler = async (req) => {
    // Write only to a private snapshot for this request.
    req.location.setWriteMode('local');

    // This change is visible only inside this handler invocation.
    req.location.root = '/var/www/' + req.getVar('user_dir');
    req.location.gzip = false;

    // ... process request ...
    // When the request pool is freed, the snapshot is discarded automatically.
};

```

6.5. Hot SSL Certificate Swap

```
// Admin handler: reload the TLS certificate from disk, zero downtime.
// New connections after this call use the new certificate;
// in-flight TLS sessions are unaffected.

adminLoc.handler = async (req) => {
  const body = JSON.parse(await req.readBody());
  const srv = nginx.http.servers.find(s => s.names.includes(body.domain));
  if (!srv) { req.respond(404, {}, 'unknown domain'); return; }

  await nginx.suspendAllWorkers();
  srv.ssl.cert = body.certPath;
  srv.ssl.key = body.keyPath;
  await nginx.resumeAllWorkers();

  req.json({ reloaded: body.domain });
};
```

7. Pre-Fork Initialization

7.1. `js_preprocess` — Config Injection During Parse

`js_preprocess` evaluates a JS script **during** `ngx_conf_parse`, at the exact point where the directive appears in the config file. The script receives a global `config` object with a single method:

`config.write(text)` — feeds `text` back into the nginx config parser exactly as if it had appeared inline at this position.

This happens **before** `init_conf`, so `nginx.http`, `nginx.stream`, and the full COM are not available. The `std` and `os` QuickJS modules are available for file and OS access.

```
nginx.conf parse reaches js_preprocess /gen.js
|
├─ New JS runtime created (isolated, short-lived)
|
├─ gen.js evaluated
|   └─ Can read files:    std.loadFile('/etc/conf.json')
|   └─ Can read env:     std.getenv('APP_ENV')
|   └─ Can run commands: os.exec([...])
|   └─ Can call config.write(text)
|
├─ Runtime freed
|
└─ Injected text re-enters the parser as if it were inline
    → full http{} / server{} / location{} blocks allowed
```

→ directives go through normal merge/optimize_servers treatment

7.1.1. Example: Server Blocks from a JSON File

```
// /etc/nginx/js/gen.js
// Generates one server {} block per tenant from a JSON config file.
// Runs once at startup (and again on `nginx -s reload`).

import * as std from 'std';

// Read tenant definitions from a JSON file.
const raw    = std.loadFile('/etc/nginx/tenants.json');
const tenants = JSON.parse(raw);

let out = '';
for (const t of tenants) {
    out += `
server {
    listen 80;
    server_name ${t.domain};
    root /var/www/${t.id};
    location / { }
}
`;
}

config.write(out);
```

```
# nginx.conf
js_preprocess /etc/nginx/js/gen.js;  # ← generates server{} blocks

events { }
http {
    # servers from gen.js appear here as if written manually
}
```

7.1.2. Example: Conditional SSL Config

```
// Enable SSL only if the certificate file actually exists.
import * as os from 'os';

const certPath = '/etc/ssl/certs/site.pem';
const keyPath  = '/etc/ssl/private/site.key';

// os.stat() returns an object on success, error code on failure.
const [stat, err] = os.stat(certPath);
if (!err) {
```

```

    config.write(`
ssl_certificate      ${certPath};
ssl_certificate_key  ${keyPath};
listen 443 ssl;
`);
}

```

7.1.3. Example: Fetching Config from a Network Service

`js_preprocess` runs synchronously inside the C parse loop. There is no event loop at this stage, so the standard `req.fetch()` API is unavailable. The `std.urlGet()` function ([QuickJS std library](#)) works around this by spawning `curl` as a child process and blocking until the response arrives.

`curl` must be installed on the system.

```

// /etc/nginx/js/fetch-config.js
// Download upstream configuration from a config service at startup.
// Uses std.urlGet() which shells out to curl – blocking and synchronous.

import * as std from 'std';

const res = std.urlGet('http://config-service.internal/v1/nginx-upstreams',
    { full: true }); // { status, headers, body }

if (res.status !== 200) {
    // Fail loudly – a bad config is worse than a failed start.
    throw new Error(`Config service returned ${res.status}`);
}

const cfg = JSON.parse(res.body);

// Generate upstream{} blocks
let out = '';
for (const up of cfg.upstreams) {
    out += `upstream ${up.name} {\n`;
    for (const peer of up.peers) {
        out += `    server ${peer.addr} weight=${peer.weight};\n`;
    }
    out += `}\n`;
}
config.write(out);

```



`std.urlGet` always spawns a `curl` subprocess. For reliability in production, prefer writing config data to a file during a pre-start hook and reading it with `std.loadFile()`, keeping `js_preprocess` independent of network availability at start time. The network fetch pattern is useful for short-lived tooling or environments where the config service is guaranteed to be reachable.

7.1.4. Multiple `js_preprocess` Directives

Multiple `js_preprocess` directives are allowed; each runs in its own isolated runtime in declaration order:

```
js_preprocess /etc/nginx/js/gen-servers.js;    # generates server{} blocks
js_preprocess /etc/nginx/js/gen-upstreams.js;  # generates upstream{} blocks
js_preprocess /etc/nginx/js/gen-maps.js;       # generates map{} blocks

events { }
http { ... }
```

7.2. `js_init_http` — Post-Parse Server Injection

`js_init_http` fires after the entire `http {}` block has been parsed but before directive merging takes place. It provides both `config.write()` and a read-only view of `nginx.http.servers[]` — the servers parsed before the `js_init_http` directive.

```
nginx.conf
├── http {
│   ├── server { server_name primary.example.com; ... }
│   ├── server { server_name secondary.example.com; ... }
│   └── js_init_http /etc/nginx/js/post-parse.js;
│       ├── nginx.http.servers[0] and [1] are readable
│       ├── config.write(text) injects new servers after these two
│       └── injected servers go through normal merge/init_locations treatment
```

```
// /etc/nginx/js/post-parse.js
// Add a management server on port 8099 only if no server is already
// listening on that port.

import * as std from 'std';

// Check whether any existing server listens on 8099
const already = nginx.http.servers.some(s =>
    s.listen.some(addr => addr.includes(':8099') || addr === '8099')
);

if (!already) {
    config.write(`
server {
    listen 8099;
```

```
server_name management.internal;

location /health { }
location /metrics { }
}
');
}
```



Place `js_init_http` **after** the `server {}` blocks it needs to read. Servers parsed after the directive are not visible to the script.

8. Workers (JS Web Workers)

JS Workers are background pthreads spawned inside an nginx worker process. They implement the standard [Web Worker API](#) and are the correct tool when a request needs to perform CPU-intensive computation without blocking the event loop.

Each nginx worker process has its own set of Workers; they are not shared between nginx workers.

8.1. Basic Worker Pattern

```
// init.js (js_source)

// Spawn a Worker thread for CPU-intensive prime computation.
// Workers are created once in js_source and reused across all requests.
const computeWorker = new Worker('/etc/nginx/js/prime_worker.js');

nginx.http.servers[0].locations
  .find(l => l.path === '/primes')
  .handler = async (req) => {
    const n = parseInt(req.queryParams.n) || 100;

    // Send work to the Worker thread and await the result.
    // While awaiting, the event loop is free to serve other requests.
    const result = await new Promise(resolve => {
      computeWorker.onmessage = e => resolve(e.data);
      computeWorker.postMessage({ n });
    });

    req.json(result);
  };
};
```

```
// /etc/nginx/js/prime_worker.js – runs in a separate pthread

function sieve(n) {
  const composite = new Uint8Array(n + 1);
```

```

    for (let i = 2; i * i <= n; i++) {
      if (!composite[i]) {
        for (let j = i * i; j <= n; j += i) composite[j] = 1;
      }
    }
    return Array.from({length: n - 1}, (_, i) => i + 2).filter(i => !composite[i]);
  }

  // Worker message loop
  onmessage = (e) => {
    const primes = sieve(e.data.n);
    postMessage({ count: primes.length, largest: primes[primes.length - 1] });
  };

```

```

HTTP request → nginx worker event loop
|
| — handler called; sends postMessage to Worker thread
|
|   □
|
|   Worker pthread
|   (independent JS runtime)
|   sieve(n) ← CPU-bound, runs without blocking event loop
|
|   □
|
|   postMessage(result)
|
|   □ onmessage event delivered to event loop
|
| — Promise resolves; req.json(result) → response sent

```

8.2. Transferring ArrayBuffers (Zero Copy)

Passing an **ArrayBuffer** by transfer moves ownership; the sender can no longer access it. Use this for large binary payloads:

```

// Compress a large buffer in a Worker thread, transfer result back.
const compressionWorker = new Worker('/etc/nginx/js/compress_worker.js');

loc.handler = async (req) => {
  const body = await req.readBody();
  const input = new TextEncoder().encode(body).buffer; // ArrayBuffer

  const compressed = await new Promise(resolve => {
    compressionWorker.onmessage = e => resolve(e.data.buf);
    // Transfer 'input' – zero copy, Worker takes ownership
    compressionWorker.postMessage({ buf: input }, [input]);
  });
};

```

```
req.setHeader('Content-Encoding', 'deflate');
req.respond(200, {}, compressed);
};
```

8.3. Worker **std** and **os** Modules

Workers have access to the QuickJS standard library:

```
// worker.js
import * as std from 'std';
import * as os from 'os';

onmessage = (e) => {
  // Blocking file read – safe in a Worker (doesn't block the event loop)
  const data = std.loadFile(e.data.path);
  postMessage({ content: data, length: data?.length ?? 0 });
};
```

See [QuickJS Standard Library](#) for the full **std/os** API reference.

9. SharedWorkers

9.1. Concept and Use Cases

A [SharedWorker](#) is a single long-lived pthread running in the nginx **master** process, connected to every nginx worker via dedicated UNIX socket channels. Unlike JS Workers (one per nginx worker), there is exactly one SharedWorker instance regardless of the **worker_processes** count.

This makes SharedWorkers the canonical location for:

- **Session stores** — a token → user mapping visible to every worker.
- **Distributed rate limiters** — accurate global request-per-second counting.
- **Feature flags** — a single flag store; updates propagate to all workers on the next message.
- **Distributed locks** — a central lock arbiter that any worker can acquire.
- **Config broadcast** — push a config change to all workers simultaneously (see [\[scope-2-all-workers-atomically\]](#)).

9.2. Port-Based Pattern (Explicit **onconnect**)

Use this pattern when a worker needs a dedicated two-way channel with the SharedWorker (i.e. replies must go back to the same worker that sent the request):

```
// /etc/nginx/js/session_sw.js – SharedWorker script
```

```

const store = new Map(); // token → {userId, expires}

// onconnect fires once per nginx worker when the channel is first used.
onconnect = (e) => {
  const port = e.ports[0];

  port.onmessage = (msg) => {
    const { cmd, token, data } = msg.data;

    if (cmd === 'get') {
      const entry = store.get(token) ?? null;
      // Reply only to the worker that asked
      port.postMessage({ token, entry });
    } else if (cmd === 'set') {
      store.set(token, { userId: data.userId,
        expires: Date.now() + data.ttl });
      port.postMessage({ ok: true });
    } else if (cmd === 'del') {
      store.delete(token);
      port.postMessage({ ok: true });
    }
  };
};

```

```

// init.js (js_source)

const sessionSW = new SharedWorker('/etc/nginx/js/session_sw.js');

// Helper: promise-based call into the SharedWorker.
function swCall(cmd, token, data) {
  return new Promise(resolve => {
    sessionSW.port.onmessage = e => resolve(e.data);
    sessionSW.port.postMessage({ cmd, token, data });
  });
}

nginx.http.servers[0].locations
  .find(l => l.path === '/auth')
  .handler = async (req) => {
    const token = req.cookies['session'];
    if (!token) { req.respond(401, {}, 'no session'); return; }

    const { entry } = await swCall('get', token, null);
    if (!entry || entry.expires < Date.now()) {
      req.respond(401, {}, 'invalid or expired session');
      return;
    }
  }

```

```
    req.setVar('authenticated_user', entry.userId);
    req.respond(200, {}, `Hello ${entry.userId}`);
};
```

9.3. Broadcast Pattern (Global **onmessage** + **clients**)

Use this pattern when the SharedWorker needs to push the same message to **all** connected workers simultaneously:

```
// /etc/nginx/js/broadcast_sw.js

// No onconnect – global onmessage handles all workers via the `clients` array.
// `clients` is an array of all currently connected port objects.

onmessage = (msg) => {
  if (msg.data.cmd === 'broadcast') {
    // Forward the message to every connected worker
    clients.forEach(c => c.postMessage(msg.data.payload));
  }
};
```

```
// init.js: trigger a broadcast from any worker's request handler.
const broadcastSW = new SharedWorker('/etc/nginx/js/broadcast_sw.js');

// Each worker registers a listener on the SW port at init time.
broadcastSW.port.onmessage = (e) => {
  // This fires on every worker when a broadcast arrives.
  const { key, value } = e.data;
  // Apply the change to this worker's local copy of the config.
  applyFeatureFlag(key, value);
};

// Admin handler: update a feature flag on all workers at once.
nginx.http.servers[0].locations
  .find(l => l.path === '/admin/feature-flag')
  .handler = async (req) => {
    const { key, value } = JSON.parse(await req.readBody());

    // Tell the SharedWorker to broadcast to all workers.
    broadcastSW.port.postMessage({
      cmd: 'broadcast',
      payload: { key, value }
    });


    req.json({ ok: true });
  };
};
```

9.4. Dynamic SharedWorker Creation

SharedWorkers created in `js_source` are spawned before `fork()` and are available to all workers from the start. They can also be created dynamically from a request handler after `fork()`:

```
// Create a SharedWorker on demand from a handler.  
// The SharedWorker is registered globally and subsequent handlers  
// calling new SharedWorker('/js/dynamic.js') reuse the same thread.  
  
loc.handler = async (req) => {  
  const sw = new SharedWorker('/etc/nginx/js/on_demand_sw.js');  
  
  const result = await new Promise(resolve => {  
    sw.port.onmessage = e => resolve(e.data);  
    sw.port.postMessage({ req: req.uri });  
  });  
  
  req.json(result);  
};
```

9.5. SharedWorker Flow Chart



10. SharedArrayBuffers and Atomics

10.1. What They Are

A `SharedArrayBuffer` is a region of memory accessible — with zero copying — from any thread: nginx workers, JS Workers, and the SharedWorker thread.

Typed array views (`Int32Array`, `Float64Array`, `BigInt64Array`, etc.) over a SAB provide direct read/write access. `Atomics` operations (`add`, `compareExchange`, `wait`, `notify`) provide the synchronisation primitives needed for safe concurrent access.

10.2. Pre-Fork SABs (Visible to All Workers)

A SAB created in `js_source` is backed by the CoW page mapping inherited by every worker after `fork()`. All workers share the same physical pages.

```
// init.js (js_source)
// A global hit counter visible to all workers with zero copy.

const counterBuf = new SharedArrayBuffer(8);
const hits       = new BigInt64Array(counterBuf);
// hits[0] is shared by every worker process.

nginx.http.servers[0].locations
  .find(l => l.path === '/count')
  .handler = (req) => {
    // Atomic increment – safe across all nginx workers simultaneously.
    const n = Atomics.add(hits, 0, 1n);
    req.text(`Request #${n + 1n}\n`);
  };

nginx.http.servers[0].locations
  .find(l => l.path === '/stats')
  .handler = (req) => {
    req.json({ total_hits: Number(Atomics.load(hits, 0)) });
  };
};
```

10.3. Post-Fork SABs (Worker-Local, Passed via Channels)

A SAB created inside a request handler (post-fork) is backed by a Linux `memfd` and can be transferred to a JS Worker or SharedWorker via `postMessage`. The receiving thread `mmap`'s the same file descriptor, establishing a shared mapping without copying data.

```
// Worker thread computes a result into a SAB; main thread reads it.
```

```

const processorWorker = new Worker('/etc/nginx/js/processor.js');

loc.handler = async (req) => {
  // Allocate shared memory for the result
  const resultBuf = new SharedArrayBuffer(1024);
  const result    = new Uint8Array(resultBuf);
  const ready     = new SharedArrayBuffer(4);
  const flag      = new Int32Array(ready);

  // Pass both SABs to the Worker; it writes the result and sets flag.
  processorWorker.postMessage({
    input:    JSON.parse(await req.readBody()),
    resultBuf, // shared – Worker writes here
    ready     // shared – Worker sets flag[0]=1 when done
  });

  // Non-blocking poll using Atomics.wait in the Worker thread
  // Meanwhile the event loop remains free.
  await new Promise(resolve => {
    processorWorker.onmessage = () => resolve();
  });

  req.respond(200, {}, new TextDecoder().decode(result));
};

```

10.4. Atomics.wait / Atomics.notify

Atomics.wait is a **blocking** call — it must only be used in Worker or SharedWorker threads, never in the main nginx event loop.

```

// Mutex implementation using SAB + Atomics.
// Safe to use in Worker and SharedWorker threads.

const mutexBuf = new SharedArrayBuffer(4); // one Int32 = one mutex
const mutex    = new Int32Array(mutexBuf);

function lock() {
  // Spin until we CAS 0 → 1 (acquire)
  while (Atomics.compareExchange(mutex, 0, 0, 1) !== 0) {
    // Block this thread until mutex[0] changes – woken by unlock()
    Atomics.wait(mutex, 0, 1);
  }
}

function unlock() {
  Atomics.store(mutex, 0, 0);
  // Wake one waiting thread
  Atomics.notify(mutex, 0, 1);
}

```

```
}
```

```
// In the SharedWorker thread – use Atomics.wait/notify to signal workers.
// SharedWorker receives work, signals the waiting nginx worker via SAB.

onconnect = (e) => {
  const port = e.ports[0];

  port.onmessage = (msg) => {
    const { sab, workItem } = msg.data;
    const signal = new Int32Array(sab);

    // Do the work ...
    const result = processWork(workItem);

    // Write result into a different part of the SAB, then notify.
    // The nginx worker is blocked in Atomics.waitAsync on this index.
    Atomics.store(signal, 0, 1);      // mark done
    Atomics.notify(signal, 0);        // wake the nginx worker's Worker thread
  };
};
```



Atomics.waitAsync (non-blocking, returns a Promise) is available in the nginx event loop context and is the correct way to wait for a SAB value to change from the main thread.

11. Accept Control and Cross-Worker Coordination

11.1. The Problem

COM mutations applied directly from a handler affect only the worker that executed the handler. Other workers have independent CoW copies and will not see the change. Furthermore, if a mutation is applied while another request on the same worker is halfway through using the old value, the result can be inconsistent.

Accept control solves both problems by temporarily pausing connection acceptance, creating a window during which no new requests can start and the worker is in a known-clean state.

11.2. Per-Worker: **suspendAcceptance** / **resumeAcceptance**

These are synchronous, instant, and affect only the calling worker.

```
loc.handler = (req) => {
```

```
// Stop accepting new connections on THIS worker.
// OS backlog holds new connections; other workers keep accepting normally.
nginx.suspendAcceptance();

// Apply a batch of mutations atomically from this worker's perspective.
const upstream = nginx.http.upstreams[0];
upstream.peers[0].weight = 10;
upstream.peers[1].weight = 5;
upstream.peers[2].down    = true;

// Resume accepting.
nginx.resumeAcceptance();

req.text('upstream updated\n');
};
```

What happens while suspended:

```
OS kernel
|
├── TCP SYN backlog accumulates (up to net.core.somaxconn connections)
|   |
|   └── When resumeAcceptance() is called, epoll fires again and
|       the worker drains the backlog in FIFO order.
```

Other workers are unaffected and continue accepting normally.

11.3. Cross-Worker: **suspendAllWorkers** / **resumeAllWorkers**

For changes that must be atomic across all workers simultaneously:

```
Worker A (handler)
|
├── await nginx.suspendAllWorkers()
|   |
|   ├── Workers B, C, ... receive coordination signal
|   |   Each finishes current request → stops accepting → ACKs
|   └── resolves when all have ACK'd
|
├── mutation window (all workers quiesced)
|
├── await nginx.resumeAllWorkers()
|   └── all workers re-add listening sockets to epoll
|
└── req.respond(...)
```

11.4. SharedWorker Broadcast + Accept Control Pattern

The most flexible pattern for cross-worker config updates:

```
// coord_sw.js – SharedWorker for cross-worker coordination

const ports = []; // one port per connected worker

onconnect = (e) => {
  const port = e.ports[0];
  ports.push(port);

  port.onmessage = (msg) => {
    if (msg.data.cmd === 'broadcast_apply') {
      // Push the same config update to every worker
      ports.forEach(p => p.postMessage({
        cmd: 'apply',
        config: msg.data.config
      }));
    }
  };
};
```

```
// init.js
const coordSW = new SharedWorker('/etc/nginx/js/coord_sw.js');

// Each worker registers a listener to apply inbound config changes.
coordSW.port.onmessage = (e) => {
  if (e.data.cmd === 'apply') {
    applyConfig(e.data.config);
    nginx.resumeAcceptance();
  }
};

function applyConfig(cfg) {
  // Apply upstream changes to this worker's local copy.
  const peers = nginx.http.upstreams[0].peers;
  cfg.peers.forEach((p, i) => {
    if (peers[i]) {
      peers[i].weight = p.weight;
      peers[i].down = p.down;
    }
  });
}

// Admin endpoint: one worker drives the update, broadcasts to all others.
nginx.http.servers[0].locations
```

```

.find(l => l.path === '/admin/config-push')
.handler = async (req) => {
  const cfg = JSON.parse(await req.readBody());

  // Quiesce all workers, apply locally, then broadcast.
  await nginx.suspendAllWorkers();

  applyConfig(cfg); // this worker applies locally

  // Tell SharedWorker to push to all workers; they will call
  // resumeAcceptance() from their onmessage handler.
  coordSW.port.postMessage({ cmd: 'broadcast_apply', config: cfg });

  // This worker also resumes acceptance.
  nginx.resumeAcceptance();

  req.json({ pushed: true });
};

```

11.5. Backlog Behaviour

While any or all workers have acceptance suspended, incoming TCP connections queue in the OS SYN/accept backlog:

- Backlog size is controlled by `net.core.somaxconn` (typically 128–4096) and the `listen` directive's `backlog=N` parameter.
- If the backlog fills before `resumeAcceptance()` is called, the kernel drops new SYN packets. From the client's perspective this looks like the server is temporarily unavailable.
- For short mutation windows (sub-millisecond), the probability of backlog overflow is negligible even under high load.



Do not hold all workers suspended for long operations (database queries, external HTTP calls). Complete all blocking work before calling `suspendAllWorkers()`, then apply changes synchronously in the mutation window.

12. Snapshots

Snapshots capture a point-in-time copy of a COM sub-tree. They are useful for rollback, for auditing config state, and for read-heavy operations that must be isolated from concurrent writes.

12.1. Basic Snapshot and Rollback

```

loc.handler = async (req) => {
  // Capture the current state of a location before mutating it.
  const snap = req.location.snapshot();

```

```

try {
  // Attempt to apply new configuration.
  const cfg = JSON.parse(await req.readBody());
  req.location.gzip      = cfg.gzip;
  req.location.root      = cfg.root;
  req.location.proxy.pass = cfg.proxyPass;

  req.json({ applied: true });
} catch (e) {
  // Something went wrong – restore the original state.
  snap.restore();
  req.respond(500, {}, `Config apply failed: ${e.message}`);
}
};

```

12.2. Sub-Object Snapshots

You can snapshot a specific COM sub-tree rather than the entire config:

```

// Snapshot just the upstream peer list.
const snap = nginx.http.upstreams[0].snapshot();

// Apply potentially risky changes.
nginx.http.upstreams[0].peers.forEach(p => p.weight = 1);

// If health-check fails, restore peer weights.
const healthy = await runHealthCheck();
if (!healthy) {
  snap.restore();
}

```

12.3. Read / Write Mode Interaction with Snapshots

When `setWriteMode("local")` is active, COM writes go to a per-request snapshot automatically, leaving the global config untouched:

```

loc.handler = async (req) => {
  // Isolate all writes to this request.
  req.location.setWriteMode('local');
  req.location.setReadMode('local');

  // Mutate the per-request copy freely.
  req.location.root = '/var/www/' + req.getVar('tenant');

  // Other concurrent requests still see the original global root.
  req.sendfile('/index.html');
}

```

```
};
```

13. Stream (TCP/UDP) Support

The JS DOM API extends to nginx's `stream {}` context. Stream handlers receive a `NginxStreamSession` object instead of `NginxRequest`.

```
js_source /etc/nginx/js/stream_init.js;

events { }
```

```
stream {
  server {
    listen 9000;
  }
}
```

```
// stream_init.js

nginx.stream.servers[0].handler = function(session) {
  // Session properties
  session.remoteAddress // client IP
  session.remotePort    // client port
  session.localAddress  // server-side IP
  session.localPort     // server-side port (9000)
  session.ssl           // boolean
  session.received      // bytes received so far

  // Variable access (stream variables)
  const proto = session.variable('protocol'); // "TCP" or "UDP"

  // Close the session with a status code
  session.finalize();           // NGX_STREAM_OK
  session.finalize(500);       // error
};
```

Stream COM mirrors the HTTP COM:

```
// Read stream upstream peers
const peers = nginx.stream.upstreams[0].peers;
peers.forEach(p => nginx.log(nginx.INFO, `${p.addr} w=${p.weight}`));

// Mutate stream upstream peers (same API as HTTP)
peers[0].weight = 20;
```

14. Bootstrap Functions (`nginx.createHttp` / `nginx.createStream`)

These two functions let you instantiate the nginx HTTP or stream module entirely from JavaScript, without an `http {}` or `stream {}` block in `nginx.conf`. They are called inside a `js_source` script at init time and may be combined with `nginx.http.addServer()`, `nginx.createSocket()`, and `nginx.http.attach()` to build a completely JS-defined server topology.

14.1. `nginx.createHttp([directives])`

Creates the HTTP module as if an empty `http {}` block were present in `nginx.conf`.

Parameter	Description
<code>directives</code> (optional)	A string containing http-context directives that would normally appear inside <code>http { ... }</code> , e.g. <code>'access_log /var/log/nginx/access.log;'</code> . Multiple directives may be separated by whitespace.

After `nginx.createHttp()` returns, `nginx.http` is available and fully functional. You can immediately call `nginx.http.addServer()`, `nginx.http.upstreams`, etc.

Full working example (from `t/hello_world.t`):

```
(function () {

    // 1. Bootstrap the HTTP module.
    nginx.createHttp('access_log /var/log/nginx/access.log;');

    // 2. Create a virtual server (nginx builds a default template when
    //     servers[] is empty).
    var srv = nginx.http.addServer('hello-world.local');

    // 3. Add a location and its JS content handler.
    srv.addLocation('/hello').handler = function (req) {
        req.respond(200, { 'content-type': 'text/plain' },
            'JS COM: Hello World!\n');
    };

    // 4. Bind a TCP socket, attach it to the HTTP listener,
    //     and register the server.
    var sock = nginx.createSocket('0.0.0.0:8080');
    nginx.http.attach(sock).addServer(srv);

})();
```

14.2. `nginx.createStream([directives])`

Creates the stream module as if an empty `stream {}` block were present in `nginx.conf`.

Parameter	Description
<code>directives</code> (optional)	A string containing stream-context directives, e.g. <code>'proxy_timeout 30s;'</code> .

After `nginx.createStream()` returns, `nginx.stream` is available.

Full working example (from `t/hello_stream.t`):

```
(function () {

    // 1. Bootstrap the stream module.
    nginx.createStream();

    // 2. Create a stream server and attach a handler.
    var srv = nginx.stream.addServer();
    srv.handler = function (session) {
        session.finalize(200);    // accept and close cleanly
    };

    // 3. Bind a TCP socket, attach it to the stream module,
    //     and register the server.
    var sock = nginx.createSocket('0.0.0.0:9000');
    nginx.stream.attach(sock).addServer(srv);

})();
```

14.3. When to Use Bootstrap Functions

Situation	Recommendation
Full server topology defined in JS (dynamic, data-driven)	Use <code>nginx.createHttp()</code> / <code>nginx.createStream()</code> . No static <code>http {}</code> / <code>stream {}</code> needed.
<code>http {}</code> block already in <code>nginx.conf</code> , only handlers vary	Skip <code>createHttp()</code> ; use <code>nginx.http.addServer()</code> and <code>addLocation()</code> to augment the existing topology.
Mix of static and dynamic servers	Include a minimal <code>http {}</code> block in <code>nginx.conf</code> for any static servers; call <code>nginx.http.addServer()</code> for additional dynamic ones.

15. Master/Worker Scripting

In addition to running handlers in worker event loops, `js_source` scripts run in the **master process** and can hook into the master's supervisory loop: reacting to worker lifecycle events, sending data to workers, and receiving data back.

15.1. Master Lifecycle Events (**nginx.on**)

Register callbacks with **nginx.on(event, fn)** in any **js_source** script. The callbacks run in the master process.

Event	Callback signature	When it fires
'workerSpawned'	fn(pid, slot)	A new worker process has started
'workerExited'	fn(pid, slot, status)	A worker process has exited
'reload'	fn()	SIGHUP received, old config still active
'reloaded'	fn()	SIGHUP complete, new workers are running
'reopen'	fn()	SIGUSR1 received (rotate logs)
'quit'	fn()	SIGQUIT received (graceful shutdown)
'message'	fn(data)	Worker sent data to master via sendToMaster() (<i>alias: 'workerMessage' with fn(slot, data)</i>)



'message' and 'workerMessage' are two spellings of the same hook. Use 'workerMessage' when you need the sender's **slot** index; use 'message' when you don't.

Example — log every worker lifecycle event to a file:

```
import * as std from 'std';

var evlog = '/var/log/nginx/worker-events.log';

function append(msg) {
    var f = std.open(evlog, 'a');
    if (f) { f.puts(msg + '\n'); f.close(); }
}

nginx.on('workerSpawned', function(pid, slot) {
    append('workerSpawned pid=' + pid + ' slot=' + slot);
});

nginx.on('workerExited', function(pid, slot, status) {
    append('workerExited pid=' + pid + ' slot=' + slot + ' status=' + status);
});

nginx.on('quit',      function() { append('quit'); });
nginx.on('reload',    function() { append('reload'); });
nginx.on('reloaded',  function() { append('reloaded'); });
nginx.on('reopen',    function() { append('reopen'); });
```

15.2. Master → Worker Messages

The master can push arbitrary JSON-serializable data to one or all workers.

Method	Description
<code>nginx.sendToWorker(slot, data)</code>	Send <code>data</code> to the worker at index <code>slot</code> . The worker receives it via <code>nginx.on('message', fn)</code> .
<code>nginx.broadcastToWorkers(data)</code>	Send <code>data</code> to every running worker. Each worker receives its own copy via <code>nginx.on('message', fn)</code> .

`data` must be JSON-serializable. It is transferred through the existing nginx channel socketpairs (the same fds used by `ngx_channel_t`), so no extra file descriptors are required.

Example — greet each worker as it spawns:

```
nginx.on('workerSpawned', function(pid, slot) {
    nginx.sendToWorker(slot, { event: 'hello', slot: slot, pid: pid });
});

// In the same js_source script, workers receive:
nginx.on('message', function(data) {
    nginx.log(nginx.INFO,
        'worker ' + nginx.workerIdx + ' received: ' + JSON.stringify(data));
});
```

15.3. Worker → Master Messages

Workers can send data back to the master at any time — typically from inside a request handler.

Method	Description
<code>nginx.sendToMaster(data)</code>	Send <code>data</code> from the current worker to the master. <code>data</code> is JSON-serializable.
<code>nginx.on('workerMessage', fn(slot, data))</code>	Register a master-side handler. <code>slot</code> is the sender's worker index; <code>data</code> is the received object.

The master's event loop is woken via SIGIO (already set up by nginx for other channel operations), so no polling is needed.

Example — worker sends a health ping; master writes it to a log file:

```
import * as std from 'std';

// Master side: receive and log
nginx.on('workerMessage', function(slot, data) {
    var f = std.open('/var/log/nginx/worker-msgs.log', 'a');
    if (f) {
```

```

        f.puts('slot:' + slot + ' msg:' + data.msg + '\n');
        f.close();
    }
});

// Worker side: send on HTTP request
nginx.http.servers[0].locations
    .find(function(l) { return l.path === '/ping'; })
    .handler = function(req) {
        nginx.sendToMaster({ slot: nginx.workerIdx, msg: 'ping' });
        req.respond(200, {'content-type': 'text/plain'}, 'sent\n');
    };

```

15.4. `nginx.workerIdx`

Inside a worker process, `nginx.workerIdx` returns the zero-based slot index of the current worker (matches the `slot` parameter in lifecycle callbacks). In the master process it returns `undefined`.

16. `nginx.repl` — Interactive JavaScript REPL

`nginx.repl` provides an in-process JavaScript REPL that evaluates expressions inside a running nginx worker over a plain HTTP upgrade connection. It is useful for live introspection, ad-hoc mutations, and debugging without a restart.

16.1. How It Works

A request handler calls `req.hijack()` to take ownership of the raw TCP connection, then calls `nginx.repl.attach()` to redirect `console.log` / `nginx.log` output to that connection, and `nginx.repl.listen()` to install a line-by-line read loop. Each line received is evaluated by `nginx.repl.eval()` and the result is written back.

16.2. REPL API

Method	Description
<code>nginx.repl.eval(line)</code>	Evaluate one line of JavaScript in the current worker's context. Returns <code>{ status, value, message, stack }</code> . <code>status</code> is "ok", "err", "incomplete", or "stack". <code>value</code> is the JSON-serialised result on success.
<code>nginx.repl.attach(fd, loglevel, nginxlevel)</code>	Start streaming <code>console.log</code> / <code>console.warn</code> / <code>console.error</code> and <code>nginx.log</code> output to <code>fd</code> . <code>loglevel</code> filters JS console methods (0 = log+warn+err, 1 = warn+err, 2 = err only); <code>nginxlevel</code> is the nginx log severity threshold.
<code>nginx.repl.detach()</code>	Stop streaming; restore normal log routing.

Method	Description
<code>nginx.repl.listen(fd, fn)</code>	Install a line-by-line read handler on <code>fd</code> . Each complete line (delimited by <code>\n</code>) is passed to <code>fn(line)</code> .
<code>nginx.repl.writeFd(fd, str)</code>	Write raw bytes to a file descriptor. Used internally to send prompt and result lines.
<code>req.hijack()</code>	Detach the HTTP connection from nginx's connection pool and return the raw file descriptor. The caller is now responsible for I/O and closing.

16.3. Wire Protocol

The client connects with an HTTP GET and the `Upgrade: nginx-repl` header. After the 101 upgrade response, the connection becomes a line-oriented plain-text stream:

```
# Client → server
<token> EVAL <javascript-expression>\n

# Server → client
<token> OK [<json-value>]\n      # success
<token> ERR <message>\n          # runtime error
<token> INCOMPLETE\n            # expression needs more input
<token> STACK <frame>\n          # stack trace line (one per frame)
```

Multiline input is handled by re-sending with accumulated lines until the server replies with something other than `INCOMPLETE`.

16.4. Embedding the REPL

```
// Install a REPL endpoint on /repl
nginx.http.servers[0].locations
  .find(function(l) { return l.path === '/repl'; })
  .handler = function(req) {
    if ((req.headers['upgrade'] || '').toLowerCase() !== 'nginx-repl') {
      return req.respond(400, {}, 'Expected Upgrade: nginx-repl\n');
    }

    var params = req.queryParams;
    var loglevel = parseInt(params.loglevel || '0');
    var nginxlevel = parseInt(params.nginxlevel || '6');

    req.respond(101, {
      'upgrade': 'nginx-repl',
      'connection': 'upgrade'
    }, '');

    var fd = req.hijack();
```

```

    nginx.repl.attach(fd, loglevel, nginxlevel);
    nginx.repl.listen(fd, function(line) {
        var r = nginx.repl.eval(line);
        if (r.status === 'ok') {
            nginx.repl._writeFd(fd,
                '- OK' + (r.value !== undefined ? ' ' + r.value : '') + '\n');
        } else if (r.status === 'incomplete') {
            nginx.repl._writeFd(fd, '- INCOMPLETE\n');
        } else {
            nginx.repl._writeFd(fd, '- ERR ' + (r.message || '') + '\n');
        }
    });
};

```

16.5. The nginx-repl Tool

A ready-made REPL server and client are provided in [js_com_tools/nginx-repl/](#):

- [nginx-repl-demo/](#) — self-contained demo: [nginx.conf](#), [repl.js](#), [repl-relay.js](#) (SharedWorker for multi-worker routing).
- [repl-client.py](#) — interactive Python client with readline support.
- [repl-client.js](#) — Node.js client.

Quick start:

```

# Build nginx with the JS module
auto/configure --add-module=src/js ...
make -j$(nproc)

# Start the demo server
./objs/nginx -c js_com_tools/nginx-repl/nginx-repl-demo/nginx.conf

# Connect with the Python client (targets worker 0 by default)
python3 js_com_tools/nginx-repl/repl-client.py

# Connect targeting a specific worker index
python3 js_com_tools/nginx-repl/repl-client.py --worker 1

```

Once connected you can evaluate any JavaScript that is valid in the worker's context:

```

> nginx.version
"1.29.7"
> nginx.workerIdx
0
> nginx.http.upstreams[0].peers.map(p => p.addr)
["192.168.1.10:8080", "192.168.1.11:8080"]
> nginx.http.upstreams[0].peers[0].weight = 10

```



The REPL evaluates JavaScript with full COM write access inside a live worker process. Bind the endpoint to loopback (**127.0.0.1**) only and protect it with a firewall or `ngx_http_access_module allow/deny` rules.

17. Applications

17.1. admin_snapshot_rollback

`js_com_apps/admin_snapshot_rollback/` is a self-contained reference application that demonstrates live configuration management without an nginx restart.

17.1.1. Concept

Operators mutate the live nginx config through the JS COM API (upstream peer weights, location handlers, added/removed servers). After a successful change, the current state is captured as a **snapshot** — a JSON ops-list that records every write made since the last base-state capture. Snapshots can be applied, rolled back, pinned, and compacted.

17.1.2. Admin HTTP API

Endpoint	Method	Description
<code>/admin/state</code>	GET	Return the current config delta vs. base state (JSON ops-list)
<code>/admin/snapshots</code>	GET	List all saved snapshots
<code>/admin/snapshots</code>	POST	Create a new snapshot from the current state
<code>/admin/snapshots/:id</code>	GET	Retrieve the JSON for snapshot id
<code>/admin/apply/:id</code>	POST	Apply snapshot id to the live config
<code>/admin/rollback</code>	POST	Step back one snapshot (undo last apply)

17.1.3. JavaScript API (`nginx.admin`)

Method	Description
<code>nginx.admin.state()</code>	Return ops-list of mutations since base-state capture
<code>nginx.admin.createSnapshot(name)</code>	Capture current state as a named snapshot
<code>nginx.admin.listSnapshots()</code>	Return array of <code>{id, name, ts}</code> objects
<code>nginx.admin.applySnapshot(id)</code>	Apply the ops-list of snapshot id
<code>nginx.admin.rollback()</code>	Revert to the state before the last <code>applySnapshot</code>
<code>nginx.admin.pin(id)</code>	Pin a snapshot as the rollback target

Method	Description
<code>nginx.admin.registerHandler(name, fn)</code>	Register a named handler function for structural ops
<code>nginx.admin.compactSnapshot(id)</code>	Reduce a snapshot to minimal ops
<code>nginx.admin.squash(ids, name)</code>	Merge multiple snapshots into one

17.1.4. `location.clearHandler()`

`location.clearHandler()` is a low-level primitive used internally by the admin application to undo handler installations during rollback.

```
// Install a handler
loc.handler = function(req) { req.respond(200, {}, 'hello'); };

// Later – remove the handler and restore the original C handler
loc.clearHandler();
```

If no JS handler was ever assigned to the location, `clearHandler()` is a no-op. After `clearHandler()`, `location.hasHandler` returns `false`.

17.1.5. Running the Demo

```
cd js_com_apps/admin_snapshot_rollback

# Start nginx (requires the JS module binary in ../../objs/nginx)
../../objs/nginx -c nginx.conf

# Inspect current state
curl http://127.0.0.1:8090/admin/state

# Create a snapshot
curl -X POST http://127.0.0.1:8090/admin/snapshots

# Apply a snapshot
curl -X POST http://127.0.0.1:8090/admin/apply/0001-initial

# Roll back
curl -X POST http://127.0.0.1:8090/admin/rollback
```

The test suite in `js_com_apps/admin_snapshot_rollback/t/` covers base-state capture, delta computation, snapshot create/apply/rollback, compaction, and structural operations (`addServer` / `addLocation` / `rollback`).

18. External Config Sources and Full Reload

18.1. Pattern A — File at Startup (**js_preprocess**)

Config data is written to a file by an external process (deploy script, Kubernetes operator, config management tool) and read synchronously during nginx startup:

```
// /etc/nginx/js/from-file.js
import * as std from 'std';

const cfg = JSON.parse(std.loadFile('/etc/nginx/dynamic-config.json') || '{}');

for (const vhost of (cfg.vhosts || [])) {
  config.write(`
server {
  listen 80;
  server_name ${vhost.domain};
  root /var/www/${vhost.docroot};
  location / { }
}
`);
}
```

This runs on every **nginx -s reload**, so the file is the single source of truth.

18.2. Pattern B — Network at Startup (**js_preprocess** + **std.urlGet**)

See the example in [\[example-fetching-config-from-a-network-service\]](#). Use this pattern for tooling environments where network availability is guaranteed. In production, pre-download to a file during a health-check or init-container step to avoid coupling nginx startup to network availability.

18.3. Pattern C — Live Reload via Handler (Zero Downtime)

For changes that do not require a full restart — upstream weights, SSL certificates, location handlers — apply mutations via the COM from a request handler:

```
// init.js

function applyConfig(cfg) {
  // Update upstream peer weights
  if (cfg.upstreams) {
    for (const [name, peerCfg] of Object.entries(cfg.upstreams)) {
      const up = nginx.http.upstreams.find(u => u.name === name);
      if (!up) continue;
      for (const [addr, w] of Object.entries(peerCfg)) {
        const peer = up.peers.find(p => p.addr === addr);
      }
    }
  }
}
```

```

        if (peer) peer.weight = w;
    }
}

// Hot-swap SSL certificates
if (cfg.ssl) {
    for (const [domain, paths] of Object.entries(cfg.ssl)) {
        const srv = nginx.http.servers.find(s => s.names.includes(domain));
        if (!srv) continue;
        srv.ssl.cert = paths.cert;
        srv.ssl.key = paths.key;
    }
}

// Admin endpoint: POST the new config JSON.
nginx.http.servers
    .find(s => s.names.includes('admin.internal'))
    .locations.find(l => l.path === '/config/apply')
    .handler = async (req) => {
        const cfg = JSON.parse(await req.readBody());

        await nginx.suspendAllWorkers();
        applyConfig(cfg);
        await nginx.resumeAllWorkers();

        req.json({ applied: true, ts: Date.now() });
    };

// Pull-based reload: a background Worker polls the config service
// periodically and applies changes.
const pollerWorker = new Worker('/etc/nginx/js/config_poller.js');
pollerWorker.onmessage = async (e) => {
    if (e.data.type === 'config-update') {
        await nginx.suspendAllWorkers();
        applyConfig(e.data.config);
        await nginx.resumeAllWorkers();
    }
};

```

```

// /etc/nginx/js/config_poller.js – JS Worker thread
import * as std from 'std';

let lastETag = '';

while (true) {
    // Poll every 30 seconds
    os.sleep(30000);
}

```

```

const res = std.urlGet('http://config-service.internal/v1/nginx-config',
                      { full: true });
const etag = (res.headers || {})[ 'etag' ] || '';

if (res.status === 200 && etag !== lastETag) {
    lastETag = etag;
    postMessage({ type: 'config-update', config: JSON.parse(res.body) });
}
}

```

18.4. Decision Guide

Requirement	js_preprocess	Pattern C (live)	Full restart
Generate server/location blocks	✓	□ (can addServer/addLocation, but no merge treatment)	✓ (via preprocess)
Network dependency at start time	Risky (prefer file)	✓	Risky
Zero downtime	□ (graceful reload)	✓	□
Upstream peer weights	✓	✓	✓
SSL certificate rotation	✓	✓	✓
New location with full directive inheritance	✓	Partial (addLocation inherits from template only)	✓

19. Third-Party Module Integration

19.1. Overview

Any nginx module can expose its configuration to the COM by implementing a small bridge file that reads the module's C config structs and projects them as JS properties.

This section walks through adding COM support to a hypothetical `ngx_http_mymod` module with a main-level config struct.

19.2. Step 1 — Identify the Config Structs

```

/* ngx_http_mymod.h */

typedef struct {
    ngx_flag_t    enable;
    ngx_uint_t    max_entries;

```

```

    ngx_str_t    backend_url;
} ngx_http_mymod_main_conf_t;

typedef struct {
    ngx_msec_t    timeout_ms;
    ngx_flag_t    cache;
} ngx_http_mymod_loc_conf_t;

extern ngx_module_t  ngx_http_mymod_module;

```

19.3. Step 2 — Write the COM Bridge File

Create `ngx_js_com_mymod.c` next to your module sources:

```

/* ngx_js_com_mymod.c */
#include <ngx_config.h>
#include <ngx_core.h>
#include <ngx_http.h>
#include "quickjs.h"
#include "ngx_js.h"
#include "ngx_http_mymod.h"

/*
 * Getter / setter for main-conf properties.
 * 'magic' selects which property (same pattern as NginxRequest).
 */
static JSValue
ngx_js_mymod_main_get(JSContext *ctx, JSValueConst this_val, int magic)
{
    ngx_http_mymod_main_conf_t *mcf;

    /* Retrieve the main config from the cycle. */
    mcf = ngx_http_cycle_get_module_main_conf(ngx_cycle,
                                                ngx_http_mymod_module);

    if (mcf == NULL) {
        return JS_UNDEFINED;
    }

    switch (magic) {
    case 0: return JS_NewBool(ctx, mcf->enable);
    case 1: return JS_NewUint32(ctx, mcf->max_entries);
    case 2: return JS_NewStringLen(ctx, (char *) mcf->backend_url.data,
                                    mcf->backend_url.len);
    }
    return JS_UNDEFINED;
}

static JSValue
ngx_js_mymod_main_set(JSContext *ctx, JSValueConst this_val,

```

```

JSValueConst val, int magic)
{
    ngx_http_mymod_main_conf_t *mcf;

    mcf = ngx_http_cycle_get_module_main_conf(ngx_cycle,
                                                ngx_http_mymod_module);

    if (mcf == NULL) {
        return JS_EXCEPTION;
    }

    switch (magic) {
    case 0:
        mcf->enable = JS_ToBool(ctx, val);
        break;
    case 1:
        JS_ToUint32(ctx, &mcf->max_entries, val);
        break;
    /* backend_url is an ngx_str_t backed by pool memory – omit setter
       or implement with careful pool allocation */
    }
    return JS_UNDEFINED;
}

/*
 * Install nginx.http.mymod into ctx.
 * Called from ngx_js_install_nginx_com() during init_conf.
 */
ngx_int_t
ngx_js_com_mymod_install(JSContext *ctx, JSValueConst nginx_obj)
{
    JSValue http_obj, mymod_obj;

    /* Retrieve the existing nginx.http object. */
    http_obj = JS_GetPropertyStr(ctx, nginx_obj, "http");
    if (JS_IsUndefined(http_obj)) {
        return NGX_ERROR;
    }

    mymod_obj = JS_NewObject(ctx);
    if (JS_IsException(mymod_obj)) {
        JS_FreeValue(ctx, http_obj);
        return NGX_ERROR;
    }

    /* Expose enable / maxEntries / backendUrl as JS properties. */
    JS_DefinePropertyGetSet(ctx, mymod_obj,
        JS_NewAtom(ctx, "enable"),
        JS_NewCFunctionMagic(ctx, ngx_js_mymod_main_get, "get_enable", 0,
                             JS_CFUNC_generic_magic, 0),
        JS_NewCFunctionMagic(ctx, ngx_js_mymod_main_set, "set_enable", 1,
                             JS_CFUNC_generic_magic, 0),

```

```

        JS_PROP_CONFIGURABLE);

    JS_DefinePropertyGetSet(ctx, mymod_obj,
        JS_NewAtom(ctx, "maxEntries"),
        JS_NewCFunctionMagic(ctx, ngx_js_mymod_main_get, "get_maxEntries", 0,
            JS_CFUNC_generic_magic, 1),
        JS_NewCFunctionMagic(ctx, ngx_js_mymod_main_set, "set_maxEntries", 1,
            JS_CFUNC_generic_magic, 1),
        JS_PROP_CONFIGURABLE);

    JS_DefinePropertyGetSet(ctx, mymod_obj,
        JS_NewAtom(ctx, "backendUrl"),
        JS_NewCFunctionMagic(ctx, ngx_js_mymod_main_get, "get_backendUrl", 0,
            JS_CFUNC_generic_magic, 2),
        JS_UNDEFINED, /* no setter – read-only */
        JS_PROP_CONFIGURABLE);

    JS_SetPropertyStr(ctx, http_obj, "mymod", mymod_obj);
    JS_FreeValue(ctx, http_obj);

    return NGX_OK;
}

```

19.4. Step 3 — Register the Bridge in the COM Initialiser

In `ngx_js_com.c`, locate `ngx_js_install_nginx_com` and add a call to your bridge function, guarded by whether the module was compiled in:

```

/* ngx_js_com.c – in ngx_js_install_nginx_com() */
#ifdef NGX_HTTP_MYMOD
    if (ngx_js_com_mymod_install(ctx, ngx_obj) != NGX_OK) {
        return NGX_ERROR;
    }
#endif

```

19.5. Step 4 — Use from JavaScript

```

// init.js
nginx.log(nginx.INFO, `mymod enabled: ${nginx.http.mymod.enable}`);
nginx.log(nginx.INFO, `mymod maxEntries: ${nginx.http.mymod.maxEntries}`);

// Toggle mymod from a handler (takes effect immediately for this worker)
nginx.http.servers[0].locations
    .find(l => l.path === '/admin/mymod-toggle')
    .handler = (req) => {

```

```

    nginx.http.mymod.enable = !nginx.http.mymod.enable;
    req.text(`mymod is now ${nginx.http.mymod.enable ? 'on' : 'off'}\n`);
};

```

19.6. Location-Level Config

For per-location config, retrieve the conf pointer from the request's `loc_conf` array rather than the cycle:

```

/* In a JS getter called from a request handler context */
static JSValue
ngx_js_mymod_loc_get_timeout(JSContext *ctx, JSValueConst this_val, int magic)
{
    ngx_http_request_t      *r;
    ngx_js_worker_t          *w;
    ngx_http_mymod_loc_conf_t *lcf;

    w = JS_GetContextOpaque(ctx);
    r = w->current_request;
    if (r == NULL) {
        return JS_UNDEFINED;
    }

    lcf = ngx_http_get_module_loc_conf(r, ngx_http_mymod_module);
    return JS_NewUint32(ctx, lcf->timeout_ms);
}

```

20. Resource Protection

20.1. Worker Memory Limit

`nginx.workerMemoryLimit` caps the QuickJS heap for all JS allocations in a worker process. The limit applies on top of nginx's own memory — it measures only memory allocated through the QuickJS allocator.

```

// Set from js_source – applies to every worker after fork().
nginx.workerMemoryLimit = 128 * 1024 * 1024; // 128 MB per worker

// Read back the configured limit (0 = unlimited).
nginx.log(nginx.INFO, `Memory limit: ${nginx.workerMemoryLimit}`);

```

When a JS allocation exceeds the limit, QuickJS throws `RangeError: out of memory`. The worker process remains healthy; the individual request receives a 500 error.

```

loc.handler = (req) => {

```

```

try {
  const huge = new ArrayBuffer(512 * 1024 * 1024); // exceeds limit
  req.text('allocated\n');
} catch (e) {
  // e is RangeError: out of memory
  req.respond(503, {}, 'memory limit exceeded, try again\n');
}
};

```

20.2. Worker Request Timeout

`nginx.workerRequestTimeout` sets the maximum wall-clock time (in milliseconds) a single JS handler invocation may consume. If the handler runs longer than this (including time spent in a tight synchronous loop), QuickJS fires an interrupt that aborts the handler. The worker survives and the request receives HTTP 500.

```

// Allow each handler at most 5 seconds of wall time.
nginx.workerRequestTimeout = 5000;

// Re-configure at runtime from a handler.
nginx.http.servers[0].locations
  .find(l => l.path === '/admin/timeout')
  .handler = (req) => {
    const ms = parseInt(req.queryParams.ms);
    if (ms >= 0) nginx.workerRequestTimeout = ms;
    req.text(`timeout is now ${nginx.workerRequestTimeout} ms\n`);
  };

```

20.3. Error Isolation

JavaScript exceptions in a handler are caught by the module framework and converted to HTTP 500 responses. The worker process never crashes due to an unhandled JS exception.

```

loc.handler = (req) => {
  // This exception is caught by the framework – worker stays up.
  throw new Error('deliberate error');
  // Client receives: HTTP/1.1 500 Internal Server Error
};

```

20.4. What JS Cannot Do

The JS runtime runs inside an nginx worker process (or master, for SharedWorkers) with the same OS privileges. The following restrictions apply:

- No access to nginx's internal C memory outside the COM API. There is no `ctypes`, `ffi`, or

equivalent.

- No subprocess spawning from handlers or SharedWorkers. The `os.exec()` function from QuickJS `os` is available only in Workers and `js_preprocess` scripts, not in nginx event-loop handlers.
- Blocking the event loop with a synchronous tight loop is interrupted by the request timeout if `nginx.workerRequestTimeout` is set.
- `Atomics.wait` (blocking) must not be called from the event loop; use `Atomics.waitAsync` (non-blocking, returns a Promise) instead.

21. Reference

21.1. Directives

Directive	Context	Arguments	Description
<code>js_source path</code>	main	1 (path)	Evaluate <code>path</code> in the master process after config parsing (<code>init_conf</code>). Multiple allowed; evaluated in declaration order. The JS runtime created here is inherited by all workers after fork.
<code>js_preprocess path</code>	main	1 (path)	Evaluate <code>path</code> immediately when encountered during <code>ngx_conf_parse</code> . Receives global <code>config.write(text)</code> . Multiple allowed; each runs in an isolated ephemeral runtime.
<code>js_init_http path</code>	http	1 (path)	Evaluate <code>path</code> after the <code>http {}</code> block finishes parsing, before merge. Receives <code>config.write(text)</code> and read-only <code>nginx.http.servers[]</code> .

21.2. `nginx` Global Object

Property / Method	Type	Description
<code>nginx.version</code>	string	nginx version string ("1.27.x")
<code>nginx.cpuCount</code>	number	Number of CPU cores
<code>nginx.workerMemoryLimit</code>	number	QuickJS heap cap in bytes; 0 = unlimited (r/w)
<code>nginx.workerRequestTimeout</code>	number	Max handler wall time in ms; 0 = unlimited (r/w)
<code>nginx.log(level, msg)</code>	—	Write to nginx error log
<code>nginx.DEBUG/INFO/NOTICE/WARN/ERR/CRIT/ALERT/EMERG</code>	number	Log level constants
<code>nginx.broadcast(msg)</code>	—	Deliver <code>msg</code> to <code>nginx.onbroadcast</code> on every worker
<code>nginx.onbroadcast</code>	function	Set handler for broadcast messages (r/w)
<code>nginx.setTimeout(ms)</code>	Promise	Resolve after <code>ms</code> ms (non-blocking)

Property / Method	Type	Description
<code>nginx.cycle</code>	object	Cycle-level stats (connections, openFilesN, sockets)
<code>nginx.events</code>	object	<code>events{}</code> directives (workerConnections, multiAccept, ...)
<code>nginx.http</code>	object	HTTP COM tree
<code>nginx.stream</code>	object	Stream COM tree
<code>nginx.suspendAcceptance()</code>	—	Remove listening sockets from this worker's epoll
<code>nginx.resumeAcceptance()</code>	—	Re-add listening sockets to this worker's epoll
<code>nginx.suspendAllWorkers()</code>	Promise	Quiesce all workers (async)
<code>nginx.resumeAllWorkers()</code>	Promise	Resume all workers (async)
<code>nginx.createHttp([directives])</code>	—	Bootstrap the HTTP module at init time (equivalent to an empty <code>http {}</code> block). <code>directives</code> is an optional string of nginx directives (e.g. <code>'access_log /var/log/nginx/access.log;'</code>). After the call, <code>nginx.http</code> is available. No-op if called more than once.
<code>nginx.createStream([directives])</code>	—	Bootstrap the stream module at init time (equivalent to an empty <code>stream {}</code> block). After the call, <code>nginx.stream</code> is available. No-op if called more than once.
<code>nginx.createSocket(addr)</code>	object	Create a new TCP listening socket bound to <code>addr</code> ("ip:port" string). Returns a socket object that can be passed to <code>nginx.http.attach(sock)</code> or <code>nginx.stream.attach(sock)</code> .
<code>nginx.on(event, fn)</code>	—	Register a master-process lifecycle callback. See [master-lifecycle-events] . Valid events: <code>'workerSpawned'</code> , <code>'workerExited'</code> , <code>'quit'</code> , <code>'reload'</code> , <code>'reloaded'</code> , <code>'reopen'</code> , <code>'message'</code> , <code>'workerMessage'</code> .
<code>nginx.sendToWorker(slot, data)</code>	—	(Master only) Send <code>data</code> to the worker at index <code>slot</code> . The worker receives it via <code>nginx.on('message', fn)</code> .
<code>nginx.broadcastToWorkers(data)</code>	—	(Master only) Send <code>data</code> to all workers. Each worker receives it via <code>nginx.on('message', fn)</code> .
<code>nginx.sendToMaster(data)</code>	—	(Worker only) Send <code>data</code> to the master process. The master receives it via <code>nginx.on('workerMessage', fn(slot, data))</code> .
<code>nginx.workerIdx</code>	number	(Worker only) Zero-based index of the current worker process (0 ... workerProcesses-1). Constant for the lifetime of the worker.
<code>nginx.repl</code>	object	Interactive JavaScript REPL object. See [nginx-repl-interactive-javascript-repl] . Properties: <code>eval(line)</code> , <code>attach(fd, loglevel, nginxlevel)</code> , <code>detach()</code> , <code>listen(fd, fn)</code> , <code>_writeFd(fd, str)</code> .

21.3. `nginx.http` Object

Property / Method	Type	Description
<code>servers[]</code>	array	Array of <code>NginxServer</code> objects
<code>upstreams[]</code>	array	Array of <code>NginxUpstream</code> objects
<code>variables</code>	object	nginx HTTP variables (r/w)
<code>serverNamesHashMaxSize</code>	number	r/w
<code>serverNamesHashBucketSize</code>	number	r/w
<code>variablesHashMaxSize</code>	number	r/w
<code>variablesHashBucketSize</code>	number	r/w
<code>sockets[]</code>	array	Listening socket objects
<code>addServer(name)</code>	<code>NginxServer</code>	Create a new virtual server at runtime
<code>removeServer(name)</code>	boolean	Remove a virtual server; returns false if not found
<code>rebuildVhostDispatch()</code>	—	Recompile the <code>server_name</code> dispatch hash
<code>match(req, conditions)</code>	boolean	Evaluate nginx <code>if</code> conditions against a request
<code>attach(sock)</code>	object	Attach a socket (from <code>nginx.createSocket()</code>) to the HTTP module; returns a listener object on which <code>addServer(srv)</code> registers a virtual server for that socket.

21.4. `NginxRequest` Methods and Properties

See [\[nginxrequest-reference\]](#).

21.5. Worker / SharedWorker API

JS DOM implements a subset of the standard Web Worker and SharedWorker APIs.

API	Notes
<code>new Worker(url)</code>	Spawn a JS Worker thread. <code>url</code> is a path to a JS file.
<code>worker.postMessage(data [, transferList])</code>	Send a value to the Worker. Transferable: <code>ArrayBuffer</code> . Shareable: <code>SharedArrayBuffer</code> .
<code>worker.onmessage = fn</code>	Receive a value from the Worker.
<code>worker.terminate()</code>	Stop the Worker thread.
<code>new SharedWorker(url)</code>	Connect to (or create) a SharedWorker. Returns an object with a <code>port</code> property.
<code>sw.port.postMessage(data)</code>	Send to the SharedWorker.
<code>sw.port.onmessage = fn</code>	Receive from the SharedWorker.

API	Notes
<code>onconnect = fn</code> (in SW script)	Fires when a new worker process connects. <code>e.ports[0]</code> is the port.
<code>onmessage = fn</code> (in SW script)	Global handler; fires for any message from any worker. <code>clients</code> array is available.
<code>clients</code> (in SW script)	Array of all connected port objects; use for broadcast.

21.6. SharedArrayBuffer and Atomics Support Matrix

Operation	Event loop	JS Worker	SharedWorker	js_preprocess
<code>new SharedArrayBuffer(n)</code>	✓	✓	✓	✓
Typed array read/write	✓	✓	✓	✓
<code>Atomics.load/store/add/...</code>	✓	✓	✓	✓
<code>Atomics.compareExchange</code>	✓	✓	✓	✓
<code>Atomics.wait</code> (blocking)	✗	✓	✓	✓
<code>Atomics.waitAsync</code>	✓	✓	✓	✗
<code>Atomics.notify</code>	✓	✓	✓	✓
Transfer via <code>postMessage</code>	✓	✓	✓	—



Never call `Atomics.wait` from the nginx event loop — it blocks the thread, preventing all request processing. Use `Atomics.waitAsync` in the event loop and `Atomics.wait` in Workers / SharedWorker threads.

21.7. Log Level Constants

Constant	nginx log level
<code>nginx.DEBUG</code>	<code>debug</code>
<code>nginx.INFO</code>	<code>info</code>
<code>nginx.NOTICE</code>	<code>notice</code>
<code>nginx.WARN</code>	<code>warn</code>
<code>nginx.ERR</code>	<code>error</code>
<code>nginx.CRIT</code>	<code>crit</code>
<code>nginx.ALERT</code>	<code>alert</code>
<code>nginx.EMERG</code>	<code>emerg</code>

22. Appendix — Flow Charts


22.1. A: Simple Synchronous Handler

Client → nginx → location matched → handler(req) → req.json() → response


22.2. B: Async Handler with Worker Offload



22.3. C: Cross-Worker COM Mutation (suspendAllWorkers)



22.4. D: SharedWorker State Coordination



22.5. E: SAB Lock-Free Counter (Pre-Fork)

js_source: `const sab = new SharedArrayBuffer(8);` ← master creates

fork() × N workers ← each inherits CoW pages
 Worker A: `Atomics.add(hits, 0, 1n)` ——— kernel mediates CAS on shared page
 Worker B: `Atomics.add(hits, 0, 1n)` ——— (no user-space lock needed)

22.6. F: Pre-Fork Config Generation

```

nginx start → conf parse → js_preprocess fires
                        |
                        JS reads files / env / network
                        |
                        config.write(generated_text)
                        |
                        text re-enters parser inline
                        |
                        server{}/location{} blocks created with
                        full merge/init_locations/optimize_servers
  
```

22.7. G: Live Reload from Config Service

```

Config service updates      Nginx workers running
    |                          |
    □                          |
    POST /config/apply        all workers quiesced
    |                          |
    □                          suspendAllWorkers()
    handler fetches / reads   applyConfig(cfg) per worker
    new config JSON           |
    |                         resumeAllWorkers()
    └─── COM mutations ────┐
                           |
                           all workers resume, serving
                           requests with new config
  
```

22.8. H: Complete nginx.* Object Tree

The following tree lists every property and method reachable from the global `nginx` object, with types and mutability. `r/o` = read-only; `r/w` = read-write; `fn` = function / method.

```

nginx (global object)
├── version          string  r/o  nginx version string, e.g. "1.27.4"
├── workerMemoryLimit number  r/w  QuickJS heap cap per worker, bytes (0 =
unlimited)
├── workerRequestTimeout number r/w  Per-handler wall-time cap, ms (0 =
unlimited)
  
```

log(level, msg)	fn	Write to nginx error log
broadcast(msg)	fn	Deliver msg to nginx.onbroadcast on every worker arrives
onbroadcast	fn	r/w Handler called when a broadcast arrives
setTimeout(ms)	Promise fn	Resolve after ms milliseconds (non-blocking)
suspendAcceptance()	fn	Remove listening sockets from this worker's epoll
resumeAcceptance()	fn	Re-add listening sockets to this worker's epoll
suspendAllWorkers()	Promise fn	Quiesce all workers (async)
resumeAllWorkers()	Promise fn	Resume all workers (async)
cycle (NginxCycle)		
hostname	string r/o	system hostname
prefix	string r/o	conf directory (dirname of nginx.conf)
confFile	string r/o	full path to nginx.conf
errorLog	string r/o	path to error log file
installPrefix	string r/o	nginx --prefix
pid	string r/o	path to PID file
connectionN	number r/o	worker_connections value
workers	number r/w	worker_processes count
daemon	bool r/o	daemon on off
master	bool r/o	master_process on off
timerResolution	number r/o	timer_resolution in ms (0 = disabled)
shutdownTimeout	number r/o	worker_shutdown_timeout in ms (0 = disabled)
priority	number r/o	worker_priority (nice value)
rlimitNofile	number r/o	worker_rlimit_nofile (-1 if unset)
workingDirectory	string r/o	working_directory (" " if unset)
sockets[]	array r/o	NginxSocket – listening sockets snapshot
events (NginxEvents)		
connections	number r/o	worker_connections
use	string r/o	event method ("epoll", "kqueue", ...)
multiAccept	bool r/w	multi_accept on off
acceptMutex	bool r/w	accept_mutex on off
acceptMutexDelay	number r/w	accept_mutex_delay in ms
http (object)		
servers[]	array r/o	NginxServer
upstreams[]	array r/o	NginxUpstream
variables	object r/o	variable name → flags/type map
serverNamesHashMaxSize	number r/o	
serverNamesHashBucketSize	number r/o	
variablesHashMaxSize	number r/o	
variablesHashBucketSize	number r/o	
sockets[]	array r/o	NginxSocket – HTTP listening sockets
addServer(name)	fn	→ NginxServer create virtual server

```

at runtime
|   |—— removeServer(name)           fn    → bool        remove virtual server
|   |—— rebuildVhostDispatch()       fn    recompile server_name dispatch hash
(call after addServer/removeServer/setNames)
|   |—— match(req, conds)            fn    → bool        evaluate nginx if-
conditions
|   |—— stream (object)
|       |—— servers[]                array  r/o      NginxStreamServer
|       |—— upstreams[]              array  r/o      NginxStreamUpstream

```

22.8.1. NginxServer properties

```

NginxServer (nginx.http.servers[i])
|
|—— name                string  r/o    primary server_name
|—— names[]             array   r/o    all server_name values
|—— listen[]            array   r/o    listen addresses
|—— locations[]         array   r/o    NginxLocation
|
|—— clientHeaderBufferSize  number r/w    bytes
|—— clientHeaderTimeout    number r/w    ms
|—— connectionPoolSize     number r/w    bytes
|—— requestPoolSize        number r/w    bytes
|—— ignoreInvalidHeaders   bool    r/w
|—— mergeSlashes           bool    r/w
|—— underscoresInHeaders   bool    r/w
|—— serverTokens           bool    r/w
|
|—— ssl                  NginxSSL (see below)
|
|—— setNames(names[])      fn
|—— addLocation(path [, {template}]) fn    → NginxLocation
|—— removeLocation(path)   fn    → bool
|—— findLocation(path)     fn    → NginxLocation | null
|—— clone(name)            fn    → NginxServer

```

22.8.2. NginxLocation properties

```

NginxLocation (server.locations[i] or server.addLocation(...))
|
|—— path                string  r/o    location path
|—— pattern              string  r/o    compiled regex (regex locations only)
|—— matchType            string  r/o    "prefix" | "exact" | "regex" | "named_regex"
|—— internal             bool    r/o
|—— jsCreated            bool    r/o    true if created by addLocation()
|
|—— handler              r/w     JS async (req) => {} content handler

```

hasHandler	bool	r/o	true if a JS handler is currently installed
clearHandler()		fn	Remove the JS handler; restore the original
nginx			
			C content handler (no-op if no handler installed)
headerFilters[]	array	r/o	installed header filter entries
bodyFilters[]	array	r/o	installed body filter entries
root	string	r/w	
alias	string	r/w	
index	string	r/w	
defaultType	string	r/w	
sendfile	bool	r/w	
tcpNopush	bool	r/w	
tcpNodelay	bool	r/w	
etag	bool	r/w	
readAhead	number	r/w	bytes
directio	number	r/w	bytes (0 = disabled)
directioAlignment	number	r/w	bytes
satisfy	string	r/w	"all" "any"
limitExcept	string	r/w	
lingering	bool	r/w	
lingeringTimeout	number	r/w	ms
lingeringTime	number	r/w	ms
keepaliveTimeout	number	r/w	ms
keepaliveRequests	number	r/w	
keepaliveTime	number	r/w	ms
keepaliveDisable	string	r/w	"none" "msie6" ...
keepaliveMinTimeout	number	r/w	ms
clientMaxBodySize	number	r/w	bytes
clientBodyTimeout	number	r/w	ms
clientBodyBufferSize	number	r/w	bytes
clientBodyInFileOnly	bool	r/w	
clientBodyInSingleBuffer	bool	r/w	
sendTimeout	number	r/w	ms
sendLowat	number	r/w	bytes
postponeOutput	number	r/w	bytes
sendfileMaxChunk	number	r/w	bytes
resetTimedoutConnection	bool	r/w	
absoluteRedirect	bool	r/w	
serverNameInRedirect	bool	r/w	
portInRedirect	bool	r/w	
chunkedTransferEncoding	bool	r/w	
resolverTimeout	number	r/w	ms
authDelay	number	r/w	ms
maxRanges	number	r/w	
ifModifiedSince	string	r/w	"off" "exact" "before"
msieRefresh	bool	r/w	
msiePadding	bool	r/w	
logNotFound	bool	r/w	
logSubrequest	bool	r/w	
recursiveErrorPages	bool	r/w	

```

|—— typesHashMaxSize  number  r/o
|
|—— ——— Module sub-objects (all r/o access; properties within are r/w) ———
|—— proxy              NginxProxy
|—— gzip               NginxGzip
|—— ssl                NginxSSL
|—— headers            NginxHeaders
|—— rewrite            NginxRewrite
|—— access             NginxAccess
|—— auth               NginxAuth
|—— authRequest         NginxAuthRequest
|—— autoindex          NginxAutoindex
|—— charset            NginxCharset
|—— dav                NginxDav
|—— fastcgi            NginxFastCGI
|—— gunzip             NginxGunzip
|—— gzipStatic         NginxGzipStatic
|—— imageFilter         NginxImageFilter
|—— limitConn          NginxLimitConn
|—— limitReq           NginxLimitReq
|—— log                NginxLog
|—— memcached          NginxMemcached
|—— mirror             NginxMirror
|—— mp4                NginxMp4
|—— proxyCache         NginxProxyCache
|—— randomIndex        NginxRandomIndex
|—— realip             NginxRealIP
|—— referer            NginxReferer
|—— scgi               NginxScgi
|—— secureLink         NginxSecureLink
|—— slice              NginxSlice
|—— ssi                NginxSsi
|—— subFilter          NginxSubFilter
|—— tryFiles           (object)
|—— userid             NginxUserid
|—— uwsgi              NginxUwsgi
|—— addition           NginxAddition
|—— xslt               NginxXslt
|
|—— setWriteMode("global"|"local"|"both")  fn
|—— setReadMode("global"|"local")          fn
|—— setProperty(name, value [, mode])      fn
|—— getProperty(name [, mode])             fn
|—— addHeaderFilter(name, fn [, priority])  fn
|—— addBodyFilter(name, fn [, priority])    fn
|—— removeHeaderFilter(name)                fn
|—— removeBodyFilter(name)                  fn
|—— getHeaderFilter(name) / getBodyFilter(name)  fn

```

22.8.3. NginxUpstream, NginxPeer, NginxRRPeer

```
NginxUpstream (nginx.http.upstreams[i] or nginx.stream.upstreams[i])
├── name          string  r/o
├── zone          string  r/o   shared memory zone name (" " if none)
├── peers[]       array   r/o   NginxPeer (config time) or NginxRRPeer (runtime)
├── addPeer(cfg)   fn     → NginxPeer | NginxRRPeer
└── removePeer(addr) fn   → bool
```

NginxPeer (config-time upstream peer)

```
├── address       string  r/o
├── weight        number  r/w
├── maxFails      number  r/w
├── failTimeout   number  r/w   seconds
├── maxConns      number  r/w
├── down         bool    r/w
└── backup        bool    r/o
```

NginxRRPeer (runtime round-robin peer – same as NginxPeer plus:)

```
├── conns        number  r/o   current active connections
├── fails        number  r/o   current fail counter
└── server       string  r/o   upstream group name
```

NginxStreamPeer / NginxStreamRRPeer – identical structure to HTTP equivalents

22.8.4. NginxSocket and NginxHttpListener / NginxStreamListener

```
NginxSocket (nginx.cycle.sockets[i] or nginx.http.sockets[i])
├── address       string  r/o   "host:port"
├── port          number  r/o
├── fd            number  r/o   OS file descriptor
├── listener      r/o      NginxHttpListener | NginxStreamListener | null
├── close()       fn     close and remove socket
└── broadcast()   fn     send fd to all workers via SCM_RIGHTS
```

NginxHttpListener (returned by socket listener attach /
nginx.http.sockets[i].listener)

```
├── address       string  r/o
├── socket        NginxSocket r/o
├── serverNames   object  r/o   { name → NginxServer }
├── serverByName(name) fn → NginxServer | null
├── addServer(cfg)   fn   → NginxServer
└── addVirtualServer(cfg) fn → NginxServer
```

NginxStreamListener – same shape as NginxHttpListener but returns NginxStreamServer

22.8.5. NginxStreamServer and NginxStreamProxy

```
NginxStreamServer (nginx.stream.servers[i])
|——— serverName          string  r/o
|——— tcpNodelay          bool    r/w
|——— prereadBufferSize   number  r/w   bytes
|——— prereadTimeout      number  r/w   ms
|——— resolverTimeout     number  r/w   ms
|——— proxyProtocolTimeout number  r/w   ms
|——— proxy               NginxStreamProxy
|——— access              NginxStreamAccess
|——— ssl                 NginxStreamSSL
|——— handler             r/w     JS (session) => {} content handler
```

```
NginxStreamProxy
|——— pass                string  r/o   upstream name
|——— connectTimeout      number  r/w   ms
|——— timeout             number  r/w   ms (proxy_timeout)
|——— nextUpstreamTimeout number  r/w   ms
|——— bufferSize          number  r/w   bytes
|——— requests            number  r/o   proxy_requests value
|——— responses           number  r/o   proxy_responses value
|——— nextUpstreamTries   number  r/w
|——— nextUpstream        string  r/w
|——— proxyProtocol       bool    r/w
|——— halfClose           bool    r/w
|——— socketKeepalive     bool    r/w
```

22.8.6. NginxRequest and NginxStreamSession

```
NginxRequest (handler argument in HTTP content handlers)
|
|——— ——— Read-only request info ———
|——— method             string   "GET" | "POST" | ...
|——— uri                string   path without query string
|——— args               string   raw query string
|——— queryParams        object   parsed {key: value} query parameters
|——— scheme             string   "http" | "https"
|——— host               string   Host header value
|——— remoteAddr         string   client IP
|——— remotePort         number   client TCP port
|——— serverAddr         string   server-side IP
|——— serverPort         number   server-side port
|——— httpVersion        string   "1.0" | "1.1" | "2.0"
|——— keepalive          bool     connection is keep-alive
|——— isInternal          bool     is subrequest
|——— startTime          number   ms since Unix epoch
|——— requestLength      number   total bytes received (headers + body)
|——— contentType        string   Content-Type header value
```

contentLength	number		Content-Length value
connection	object		nginx connection info
headers	object		request headers (lowercased keys)
cookies	object		parsed cookies {name: value}
upstream	object		upstream peer info (if proxied)
——— Mutable ———			
statusCode	number	r/w	response status code
variables	object	r/w	nginx variables exotic proxy (r.variables.foo)
location	NginxLocation	r/o	matched location (mutable via setWriteMode)
——— Response ———			
respond(status, headers, body)	fn		send complete response
json(data)	fn		200 + application/json
text(body)	fn		200 + text/plain
html(body)	fn		200 + text/html
redirect(url)	fn		302 redirect
writeHead(status, headers)	fn		begin streaming response
write(chunk)	fn		append streaming chunk
finish()	fn		end streaming response
sendfile(path)	fn		serve file via sendfile(2)
——— Headers ———			
setHeader(name, value)	fn		
getHeader(name)	fn		→ string undefined
removeHeader(name)	fn		
——— Variables ———			
variable(name) / getVar(name)	fn		→ string number undefined
setVariable(name, v) / setVar	fn		
——— Body ———			
body	string	r/o	buffered body (null if not yet read)
readBody()	Promise	fn	read full request body → string
——— Async ———			
sleep(ms)	Promise	fn	non-blocking delay
fetch(url)	Promise	fn	non-blocking upstream HTTP request → Response
subrequest(uri)	Promise	fn	nginx internal subrequest → Response
log(level, msg)	fn		

NginxStreamSession (handler argument in stream content handlers)

remoteAddress	string	r/o	
remotePort	number	r/o	
localAddress	string	r/o	
localPort	number	r/o	
received	number	r/o	bytes received so far
status	number	r/o	session status code
ssl	object	r/o	SSL info (null on plain TCP)
finalize([code])	fn		close session; default code = NGX_STREAM_OK

└── variable(name) fn → string | undefined

22.8.7. Worker and SharedWorker

Worker (new Worker(url))

└── onmessage fn r/w called when the Worker posts a message
└── postMessage(data [, transferList]) fn send value to the Worker
└── terminate() fn stop the Worker thread

SharedWorker (new SharedWorker(url))

└── port MessagePort

MessagePort (sw.port)

└── onmessage fn r/w called when the SharedWorker posts a message
└── postMessage(data) fn send value to the SharedWorker

—— Inside a SharedWorker script ——

onconnect fn r/w fires when a new nginx worker connects;
e.data.ports[0] = MessagePort
onmessage fn r/w global handler; receives messages from any worker
clients[] array all connected MessagePort objects (for broadcast)

22.8.8. NginxSSL and NginxStreamSSL

NginxSSL (server.ssl in http context)

└── sessionTimeout number r/w seconds
└── sessionTickets bool r/w
└── preferServerCiphers bool r/w
└── verifyDepth number r/w
└── ciphers string r/o
└── clientCertificate string r/o
└── trustedCertificate string r/o
└── ecdhCurve string r/o
└── dhparam string r/o
└── setCiphers(list) fn change cipher list
└── setProtocols(list) fn e.g. "TLSv1.2 TLSv1.3"
└── setCertificate(cert, key) fn hot-swap certificate and key

NginxStreamSSL (streamServer.ssl in stream context) – same properties plus:

└── handshakeTimeout number r/w ms

22.9. I: Full Nginx JS Class Reference

22.9.1. Class Index

Class name	JS constructor / access path	Purpose
<code>NginxCycle</code>	<code>nginx.cycle</code>	Cycle-level nginx globals
<code>NginxEvents</code>	<code>nginx.events</code>	<code>events{}</code> block settings
<code>NginxServer</code>	<code>nginx.http.servers[i]</code>	HTTP virtual server
<code>NginxLocation</code>	<code>server.locations[i], addLocation()</code>	HTTP location block
<code>NginxUpstream</code>	<code>nginx.http.upstreams[i]</code>	HTTP upstream group
<code>NginxPeer</code>	<code>upstream.peers[i]</code> (config time)	Upstream server (static)
<code>NginxRRPeer</code>	<code>upstream.peers[i]</code> (runtime)	Upstream server (live stats)
<code>NginxRequest</code>	handler argument <code>req</code>	HTTP request/response
<code>NginxSocket</code>	<code>nginx.cycle.sockets[i], http.sockets[i]</code>	Bound listening socket fd
<code>NginxHttpListener</code>	<code>socket.listener</code>	HTTP virtual-host router
<code>NginxStreamServer</code>	<code>nginx.stream.servers[i]</code>	Stream (TCP/UDP) server
<code>NginxStreamProxy</code>	<code>streamServer.proxy</code>	Stream proxy settings
<code>NginxStreamListener</code>	<code>socket.listener (stream)</code>	Stream virtual-host router
<code>NginxStreamSession</code>	handler argument <code>session</code>	TCP/UDP session
<code>NginxStreamUpstream</code>	<code>nginx.stream.upstreams[i]</code>	Stream upstream group
<code>NginxStreamPeer</code>	<code>streamUpstream.peers[i]</code> (config)	Stream upstream server (static)
<code>NginxStreamRRPeer</code>	<code>streamUpstream.peers[i]</code> (runtime)	Stream upstream server (live)
<code>Worker</code>	<code>new Worker(url)</code>	Background JS thread (per nginx-worker)

Class name	JS constructor / access path	Purpose
SharedWorker	<code>new SharedWorker(url)</code>	Shared JS thread (one, in master)
MessagePort	<code>sharedWorker.port</code>	Channel endpoint to SharedWorker
NginxSSL	<code>server.ssl</code>	HTTP SSL/TLS settings
NginxStreamSSL	<code>streamServer.ssl</code>	Stream SSL/TLS settings
NginxProxy	<code>location.proxy</code>	<code>proxy_pass</code> and proxy settings
NginxGzip	<code>location.gzip</code>	<code>gzip</code> module settings
NginxHeaders	<code>location.headers</code>	<code>add_header</code> / <code>expires</code> settings
NginxRewrite	<code>location.rewrite</code>	<code>rewrite</code> module settings
NginxAccess	<code>location.access</code>	<code>allow/deny</code> access rules
NginxAuth	<code>location.auth</code>	<code>auth_basic</code> settings
NginxAuthRequest	<code>location.authRequest</code>	<code>auth_request</code> settings
NginxAutoindex	<code>location.autoindex</code>	Directory listing settings
NginxCharset	<code>location.charset</code>	Charset conversion settings
NginxDav	<code>location.dav</code>	WebDAV settings
NginxFastCGI	<code>location.fastcgi</code>	FastCGI proxy settings
NginxGunzip	<code>location.gunzip</code>	On-the-fly decompression
NginxGzipStatic	<code>location.gzipStatic</code>	Pre-compressed file serving
NginxImageFilter	<code>location.imageFilter</code>	Image transformation settings
NginxLimitConn	<code>location.limitConn</code>	<code>limit_conn</code> settings
NginxLimitReq	<code>location.limitReq</code>	<code>limit_req</code> settings
NginxLog	<code>location.log</code>	Access log settings
NginxMemcached	<code>location.memcached</code>	Memcached proxy settings
NginxMirror	<code>location.mirror</code>	Request mirroring settings
NginxMp4	<code>location.mp4</code>	MP4 pseudo-streaming settings

Class name	JS constructor / access path	Purpose
NginxProxyCache	location.proxyCache	Proxy cache settings
NginxRandomIndex	location.randomIndex	Random index file settings
NginxRealIP	location.realIp	Real-IP header settings
NginxReferer	location.referer	Referer validation settings
NginxScgi	location.scgi	SCGI proxy settings
NginxSecureLink	location.secureLink	Secure link settings
NginxSlice	location.slice	Response slicing settings
NginxSsi	location.ssi	SSI processing settings
NginxSubFilter	location.subFilter	Response substitution settings
NginxUserid	location.userId	User-ID cookie settings
NginxUwsgi	location.uwsgi	uWSGI proxy settings
NginxAddition	location.addition	Response addition settings
NginxXslt	location.xslt	XSLT transformation settings
NginxStreamAccess	streamServer.access	Stream access rules

22.9.2. Module COM Class Properties

Class	Property	Mode	Notes
NginxProxy	pass	r/o	upstream URL or name
NginxProxy	httpVersion	r/w	"1.0" or "1.1"
NginxProxy	connectTimeout	r/w	ms
NginxProxy	sendTimeout	r/w	ms
NginxProxy	readTimeout	r/w	ms
NginxProxy	buffering	r/w	bool
NginxProxy	requestBuffering	r/w	bool
NginxProxy	interceptErrors	r/w	bool
NginxProxy	bufferSize	r/w	bytes
NginxProxy	nextUpstreamTries	r/w	number
NginxProxy	nextUpstreamTimeout	r/w	ms

Class	Property	Mode	Notes
NginxGzip	enable	r/w	bool
NginxGzip	level	r/w	1–9
NginxGzip	minLength	r/w	bytes
NginxGzip	vary	r/w	bool — add Vary: Accept-Encoding
NginxHeaders	expires	r/o	seconds; -1 = epoch; -2 = max; 0 = off
NginxHeaders	expiresTime	r/o	absolute Unix timestamp
NginxHeaders	headersInherit	r/w	bool
NginxHeaders	trailersInherit	r/w	bool
NginxRewrite	log	r/w	bool — log rewrite rule execution
NginxRewrite	uninitializedVariableWarn	r/w	bool
NginxRewrite	stackSize	r/w	number — rewrite stack depth
NginxRewrite	hasRules	r/o	bool — any rules defined
NginxAuth	realm	r/w	string — Basic auth realm
NginxAuth	userFile	r/w	string — path to htpasswd file
NginxAuthRequest	uri	r/w	string — subrequest URI
NginxAutoindex	enable	r/w	bool
NginxAutoindex	format	r/w	"html" "xml" "json" "jsonp"
NginxAutoindex	localtime	r/w	bool
NginxAutoindex	exactSize	r/w	bool
NginxCharset	charset	r/o	string
NginxCharset	sourceCharset	r/o	string
NginxCharset	overrideCharset	r/w	bool
NginxCharset	setCharset(from, to)	fn	add charset conversion mapping
NginxDav	methods	r/w	string[] — allowed HTTP methods
NginxDav	access	r/w	number — octal permissions (e.g. 0o644)
NginxDav	minDeleteDepth	r/w	number
NginxDav	createFullPutPath	r/w	bool
NginxFastCGI	pass	r/o	string
NginxFastCGI	index	r/w	string
NginxFastCGI	keepConn	r/w	bool
NginxFastCGI	connectTimeout	r/w	ms
NginxFastCGI	sendTimeout	r/w	ms

Class	Property	Mode	Notes
NginxFastCGI	readTimeout	r/w	ms
NginxFastCGI	buffering	r/w	bool
NginxFastCGI	requestBuffering	r/w	bool
NginxFastCGI	interceptErrors	r/w	bool
NginxGunzip	enable	r/w	bool
NginxGzipStatic	enable	r/w	bool — serve pre-compressed .gz files
NginxImageFilter	action	r/w	"test" "resize" "crop" "rotate"
NginxImageFilter	width	r/w	pixels
NginxImageFilter	height	r/w	pixels
NginxImageFilter	angle	r/w	90 180 270
NginxImageFilter	jpegQuality	r/w	1–100
NginxImageFilter	webpQuality	r/w	1–100
NginxImageFilter	sharpen	r/w	0–100
NginxImageFilter	transparency	r/w	bool
NginxImageFilter	interlace	r/w	bool
NginxImageFilter	bufferSize	r/w	bytes
NginxLimitReq	logLevel	r/w	log level string
NginxLimitReq	delayLogLevel	r/w	log level string
NginxLimitReq	statusCode	r/w	HTTP status for rejected requests (default 503)
NginxLimitReq	dryRun	r/w	bool — count but do not limit
NginxLimitConn	logLevel	r/w	log level string
NginxLimitConn	statusCode	r/w	HTTP status (default 503)
NginxLimitConn	dryRun	r/w	bool
NginxLog	off	r/w	bool — disable access logging for this location
NginxMemcached	connectTimeout	r/w	ms
NginxMemcached	sendTimeout	r/w	ms
NginxMemcached	readTimeout	r/w	ms
NginxMemcached	bufferSize	r/w	bytes
NginxMemcached	gzipFlag	r/w	number — Memcached flags value for gzip content
NginxMirror	uris	r/w	string[] — mirror target URIs

Class	Property	Mode	Notes
NginxMirror	requestBody	r/w	bool — mirror request body
NginxMp4	bufferSize	r/w	bytes
NginxMp4	maxBufferSize	r/w	bytes
NginxMp4	startKeyFrame	r/w	bool
NginxRandomIndex	enable	r/w	bool
NginxRealIP	header	r/o	string — header used as client IP source
NginxRealIP	recursive	r/w	bool
NginxReferer	noReferer	r/w	bool — allow requests without Referer
NginxReferer	blockedReferer	r/w	bool — allow requests with blocked Referer
NginxReferer	serverNames	r/w	string[] — allowed referer server names
NginxScgi	connectTimeout	r/w	ms
NginxScgi	sendTimeout	r/w	ms
NginxScgi	readTimeout	r/w	ms
NginxScgi	bufferSize	r/w	bytes
NginxUwsgi	connectTimeout	r/w	ms
NginxUwsgi	sendTimeout	r/w	ms
NginxUwsgi	readTimeout	r/w	ms
NginxUwsgi	bufferSize	r/w	bytes
NginxSecureLink	secret	r/w	string — HMAC secret
NginxSecureLink	variable	r/o	string — variable name holding the link
NginxSecureLink	md5	r/o	bool — MD5-based (vs HMAC-SHA1)
NginxSlice	size	r/w	bytes — slice size for range-splitting
NginxSsi	enable	r/w	bool
NginxSsi	silentErrors	r/w	bool
NginxSsi	lastModified	r/w	bool
NginxSsi	valueLen	r/w	number — max SSI variable value length
NginxSubFilter	once	r/w	bool — replace only first occurrence
NginxSubFilter	lastModified	r/w	bool
NginxUserid	enable	r/w	bool "v1" "v2" "log"

Class	Property	Mode	Notes
NginxUserid	name	r/w	string — cookie name
NginxUserid	domain	r/w	string — cookie domain
NginxUserid	path	r/w	string — cookie path
NginxUserid	p3p	r/w	string — P3P header value
NginxUserid	expires	r/o	number — seconds
NginxUserid	mark	r/w	string — version mark character
NginxAddition	addBeforeBody	r/w	string — URI to prepend to response body
NginxAddition	addAfterBody	r/w	string — URI to append to response body
NginxProxyCache	minUses	r/w	number
NginxProxyCache	lock	r/w	bool
NginxProxyCache	lockTimeout	r/w	ms
NginxProxyCache	lockAge	r/w	ms
NginxProxyCache	revalidate	r/w	bool
NginxProxyCache	convertHead	r/w	bool
NginxProxyCache	backgroundUpdate	r/w	bool
NginxXslt	lastModified	r/w	bool
NginxXslt	sheetsCount	r/o	number — number of loaded stylesheets